

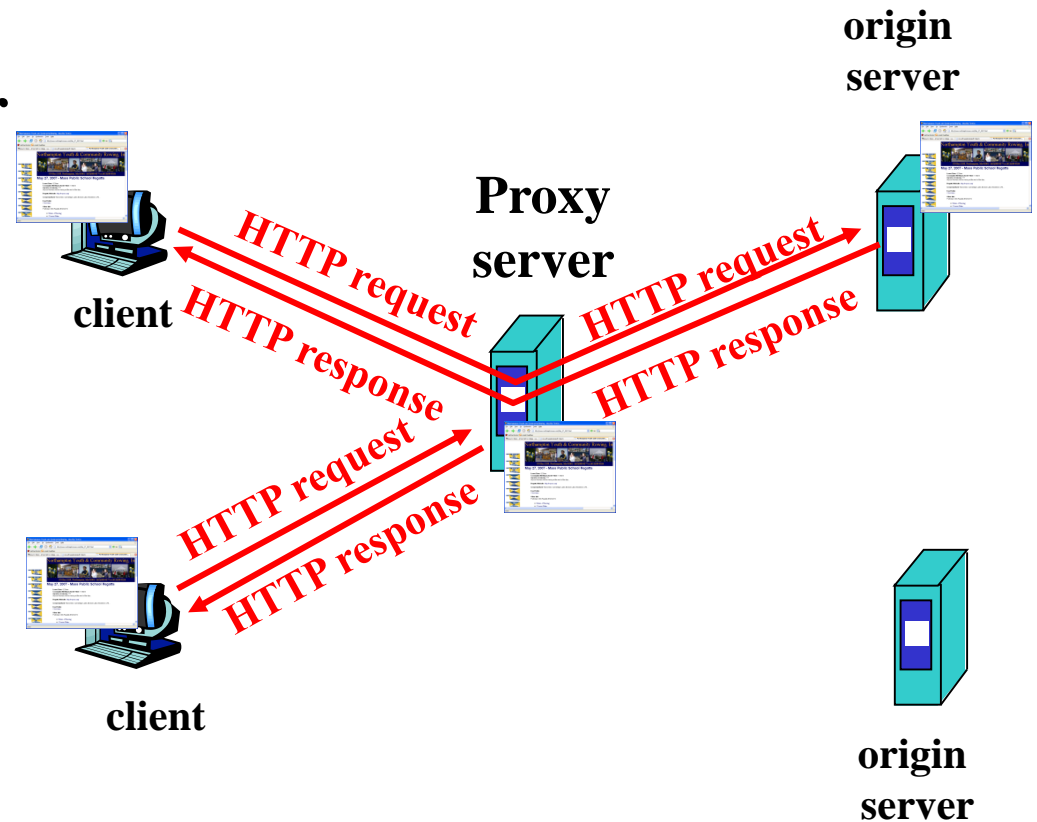
Content Distribution Networks and Peer-to-Peer

**Mestrado Integrado em
Engenharia de Computadores e
Telemática
2019/2020**

Still from FR: Web caches (proxy server)

Goal: satisfy client request without involving origin server

- user sets browser: Web accesses via proxy server
- browser sends all HTTP requests to proxy
 - object in cache: cache returns object
 - else proxy requests object from origin server, then returns object to client



More about Web caching

- Proxy server acts as both client and server
- typically proxy server is installed by ISP (university, company, residential ISP)

Why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link.


Caching example

Assumptions

- average object size = 100,000 bits
- avg. request rate from institution's browsers to origin servers = 15/sec
- delay from institutional router to any origin server and back to router = 2 sec

Consequences

- utilization on LAN = 15%
- utilization on access link = 100%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + milliseconds




Caching example (cont)

possible solution

- increase bandwidth of access link to, say, 10 Mbps

consequence

- utilization on LAN = 15%
- utilization on access link = 15%
- Total delay = Internet delay + access delay + LAN delay
= 2 sec + msecs + msecs
- often a costly upgrade




Caching example (cont)

possible solution: install cache

- suppose hit rate is 0.4

consequence

- 40% requests will be satisfied almost immediately
- 60% requests satisfied by origin server
- utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- total avg delay = Internet delay + access delay + LAN delay = $.6 \cdot (2.01) \text{ secs} + .4 \cdot \text{milliseconds} < 1.4 \text{ secs}$




Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
- **cache:** specify date of cached copy in HTTP request
If-modified-since:
<date>
- **server:** response contains no object if cached copy is up-to-date:

HTTP/1.0 304 Not Modified

cache

server




CDNs vs Web Caches

- **Caches are used by ISPs to reduce bandwidth consumption**
- **CDNs are used by content providers to improve quality of service to end users**
- **Caches are reactive**
- **CDNs are proactive**
- **Caching proxies cater to their users (web clients) and not to content providers (web servers)**
- **CDNs cater to the content providers (web servers) and clients**
- **CDNs give control over the content to the content providers**
- **Caching proxies do not**

So much happened in our digitalized world in 2017 – and we have the numbers behind it

Things that happened online in 2017 within 60 seconds

Lots of multimedia content!




Motivation (1)

- **IP based networks**
- **Web based applications have become the norm for corporate internal networks and many business-to-business interactions**
- **Large acceptance and explosive growth**
 - **Serious performance problems**
 - **Degraded user experience**
 - For a large set of applications, including VIDEO access**
- **Improving the performance of networked applications**
 - **Use many sites at different points within the network**
 - Stand alone servers
 - Routers

Motivation (2): Flash Crowds:

Consider, On September 11, 2001



Content Distribution Network

What is a CDN?

A network of servers delivering content on behalf of an origin site

Large-file service with

- No custom client**
- No custom server**
- No prepositioning**

Content distribution networks

- Client attempts to access the main server site for an application
- It is redirected to one of the other sites
- Each site caches information
 - Avoid going to the main server to get the information/application
- Access a closely located site
 - Avoid congestion on the path to the main server
- Set of sites used to improve the performance of web-based applications collectively
 - Content distribution network

CDNs basics

- **A number of CDN companies well established now**
 - E.g. Akamai, Digital Island, Speedera, CDN77, Cloudfare, Stackpat
- **Many companies are exploring CDNs**
 - **Avoid congested portions of the Internet**
 - E.g., CNN, CNBC, ...
- **Consist of**
 - **Edge servers deployed at several ISP (Internet Service Provider) access locations and network exchange points**
- **Improve the response time of an Internet site**
 - **Offloading the delivery of bandwidth-intensive objects, such as images and video clips**
- **Intelligent Internet infrastructure that improves the performance and scalability of distributed applications by moving the bulk of their *computation* to servers located at the edge of the network**

CDN challenges (I)

- **Keep consistency among the enterprise data hosted by the offloaded applications**
- **Share session state among edge and origin application servers**
- **Distribution, configuration, and management**
- **Application security.**

CDN challenges (II)

- **Load-sharing content**
 - **Handle requests fairly amongst servers/sites**
 - **Easily add servers/sites to content service**
 - **Adjust connections based on server/site load**
- **Content Availability with multiple servers**
 - **Synchronize content amongst servers/sites**
 - **Avoid faulty servers/sites**
- **Handle applications which use ‘state’**
 - **Need to learn client ID to satisfy state requirement**
 - **Need to maintain state for period of time - variable**

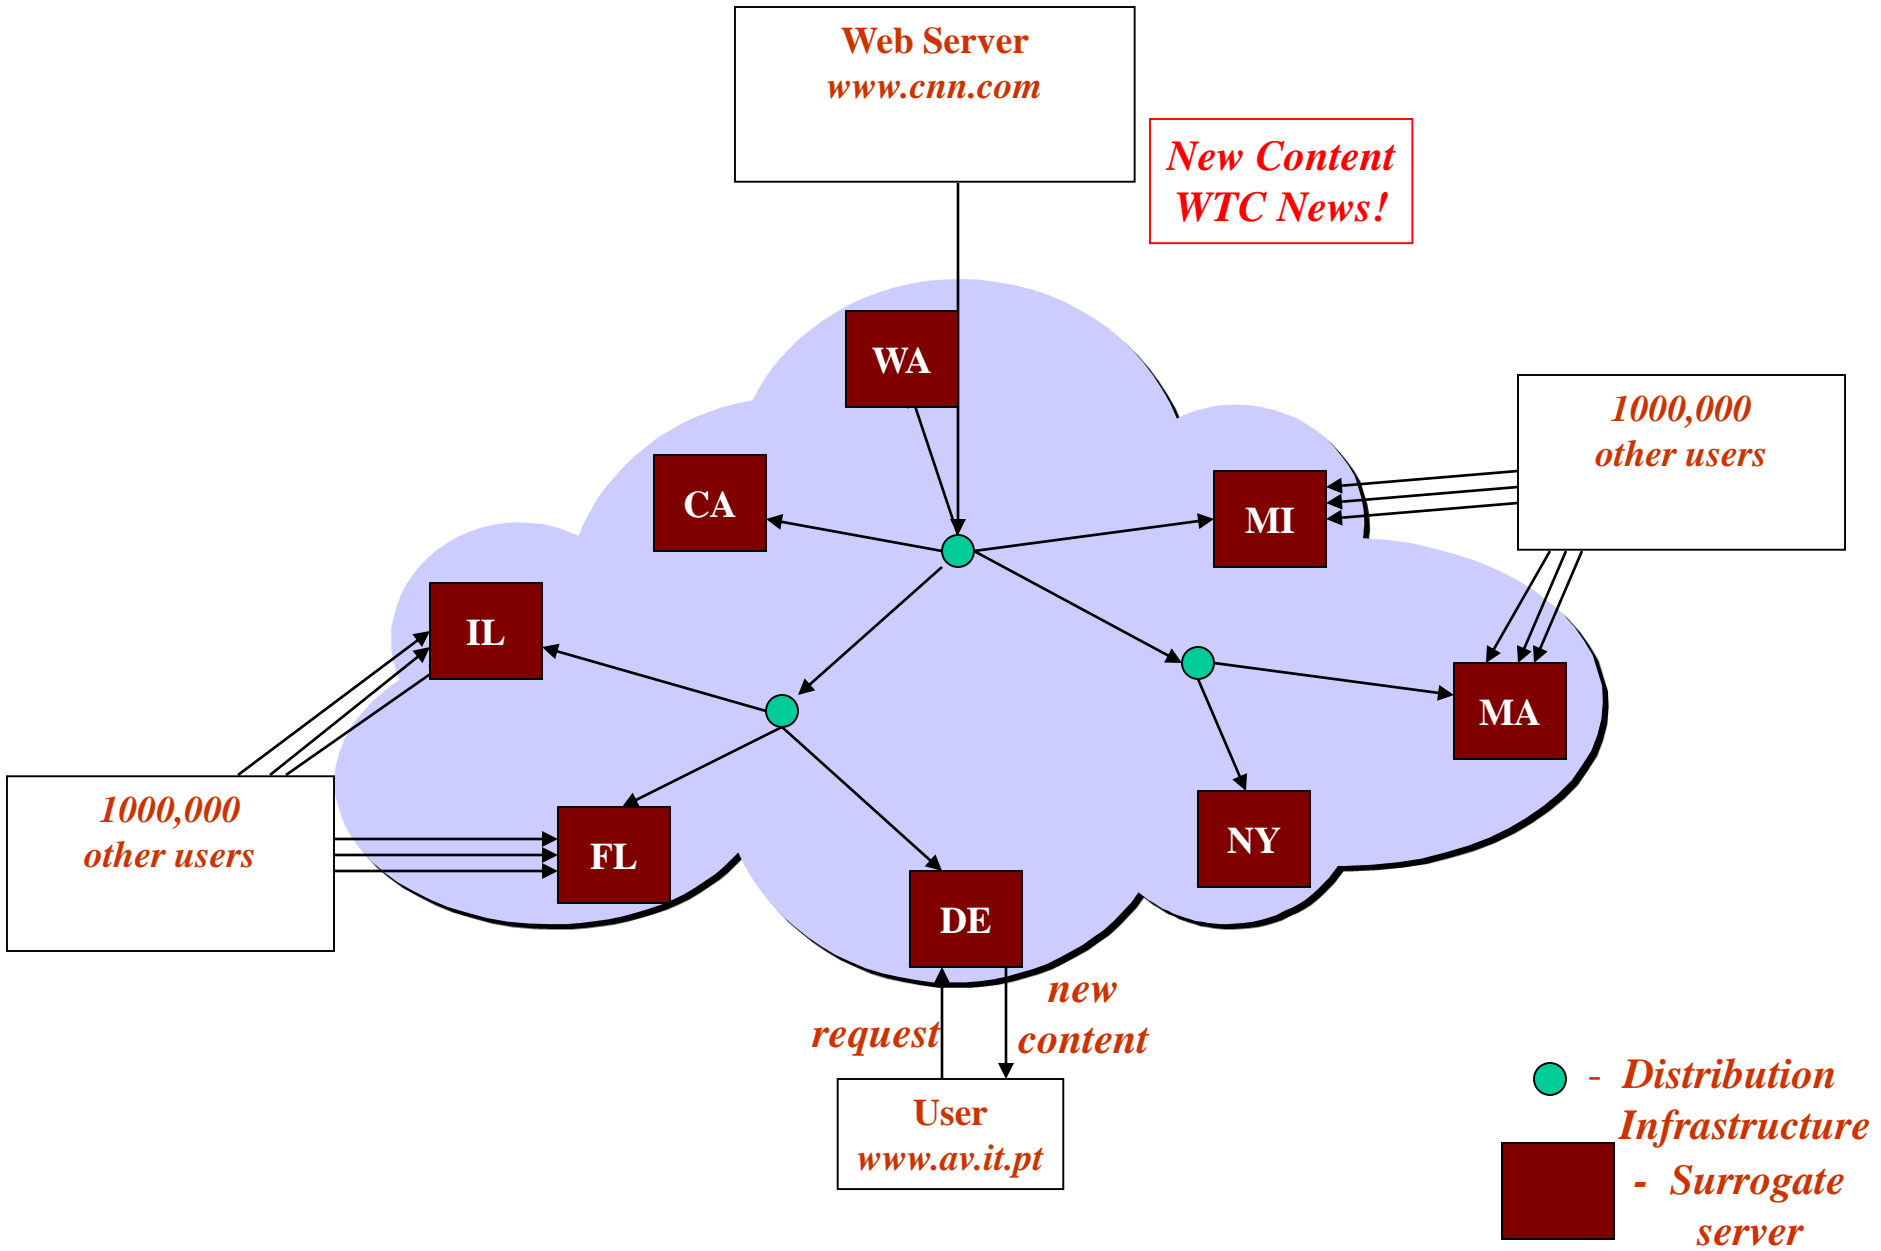
What is a CDN?

- **Content Delivery Network**
 - **Also sometimes called Content Distribution Network**
 - **At least half of the world's bits are delivered by a CDN**
 - Probably closer to 80/90%
- **Primary Goals**
 - **Create replicas of content throughout the Internet**
 - **Ensure that replicas are always available**
 - **Direct clients to replicas that will give good performance**

Key Components of a CDN

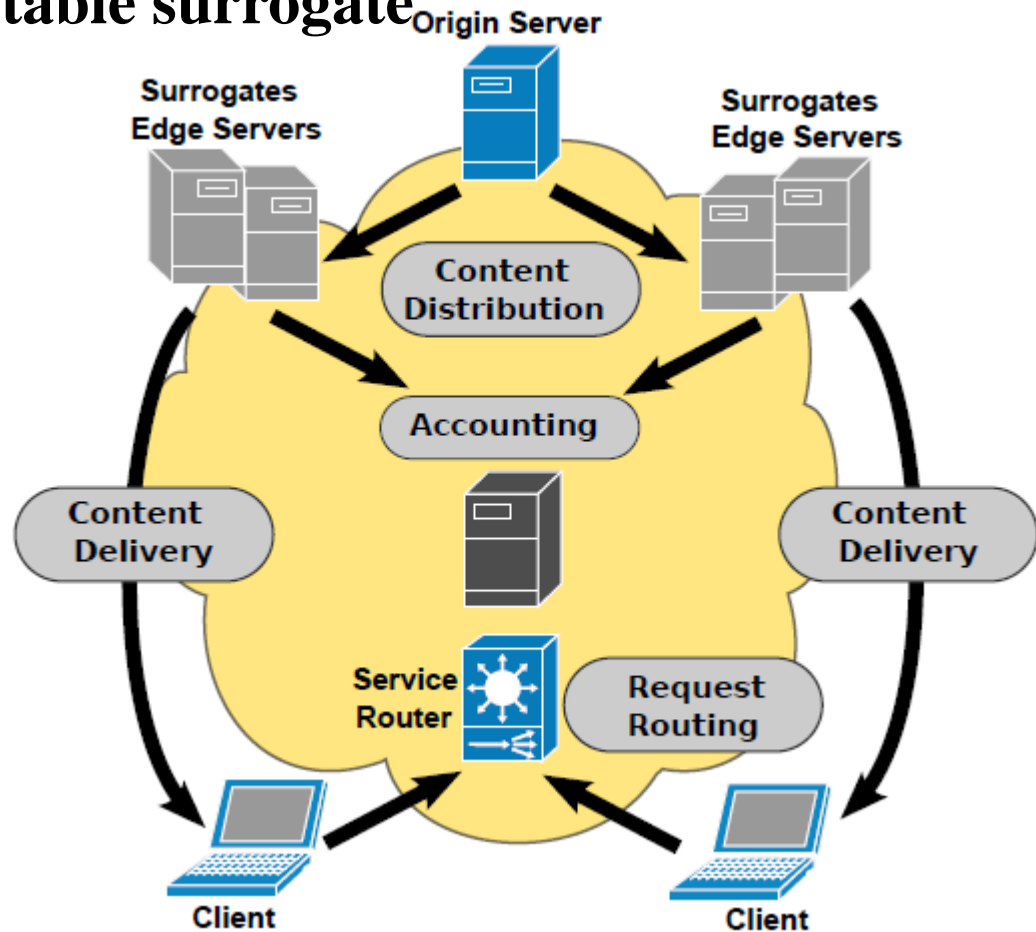
- **Distributed servers**
 - Usually located inside of other ISPs
- **High-speed network connecting them**
- **Clients**
 - Can be located anywhere in the world
 - They want fast Web performance
- **Binding clients and distributed servers**
 - Something that binds clients to “nearby” replica servers

CDN Architecture



CDN Components

- ***Content Delivery Infrastructure:*** Delivering content to clients from surrogates
- ***Request Routing Infrastructure:*** Steering or directing content request from a client to a suitable surrogate
- ***Distribution Infrastructure:*** Moving or replicating content from content source (origin server, content provider) to surrogates
- ***Accounting Infrastructure:*** Logging and reporting of distribution and delivery activities



Inside a CDN

- **Servers are deployed in clusters for reliability**
 - **Some may be offline**
 - Could be due to failure
 - Also could be “suspended” (e.g., to save power or for upgrade)
- **Could be multiple clusters per location (e.g., in multiple racks)**
- **Server locations**
 - **Well-connected points of presence (PoPs)**
 - **Inside of ISPs**

Mapping clients to servers

- **CDNs need a way to send clients to the “best” server**
 - The best server can change over time
 - And this depends on client location, network conditions, server load, ...
 - What existing technology can we use for this?
- **DNS-based redirection**
 - Clients request www.foo.com
 - DNS server directs client to one or more IPs based on request IP
 - Use short TTL to limit the effect of caching

CDN redirection example

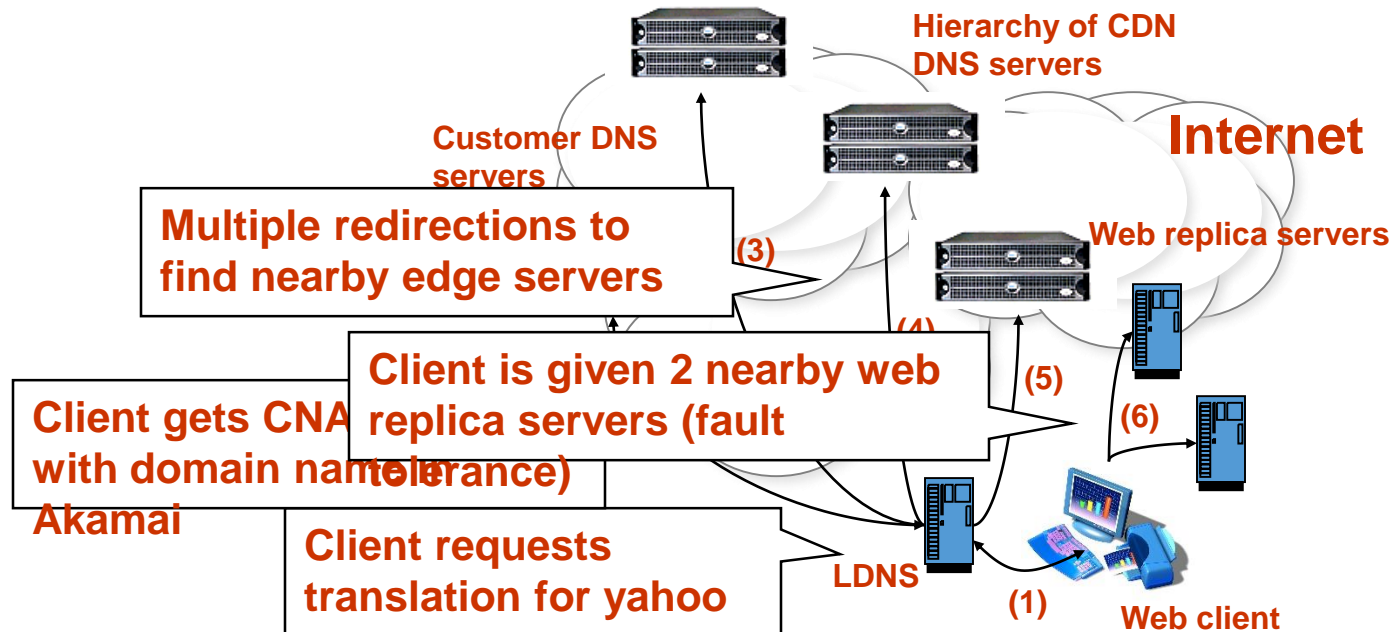
```
choffnes$ dig www.fox.com
```

```
;; ANSWER SECTION:
```

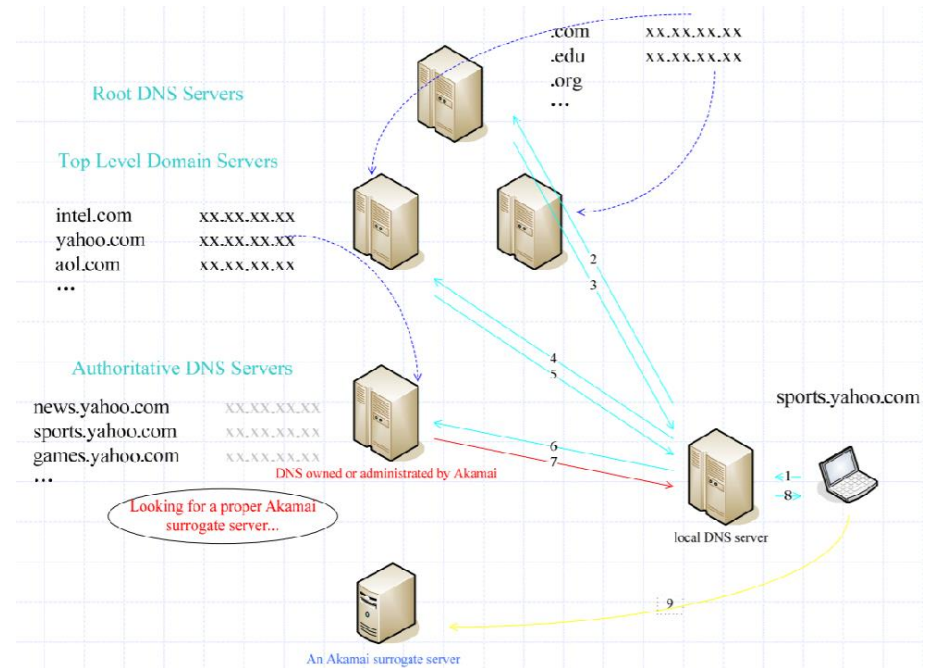
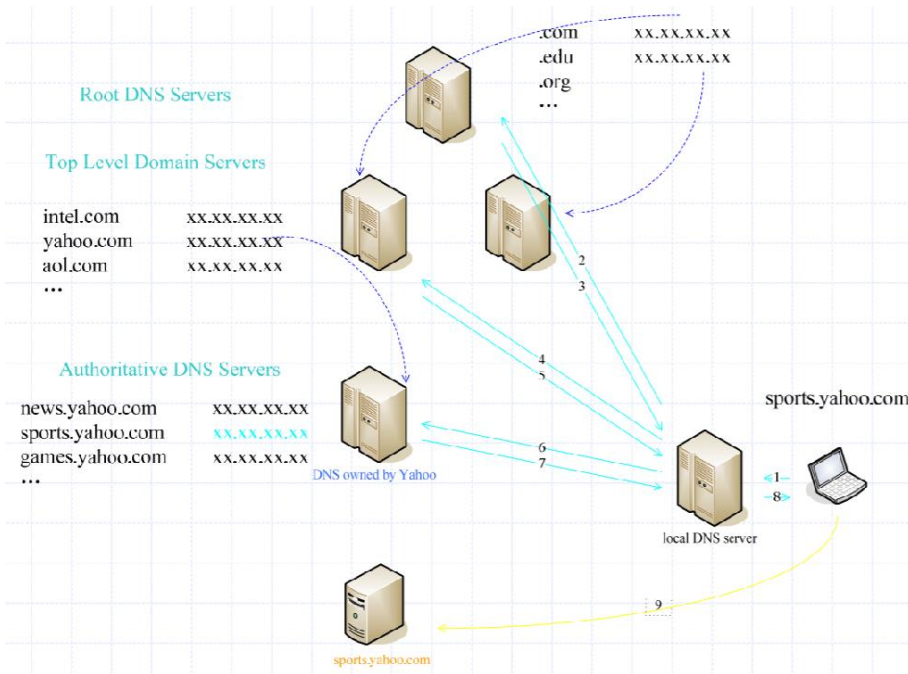
www.fox.com.	510	IN	CNAME	www.fox-
rma.com.edgesuite.net.				
www.fox-rma.com.edgesuite.net.	5139	IN	CNAME	a2047.w7.akamai.net.
a2047.w7.akamai.net.	4	IN	A	23.62.96.128
a2047.w7.akamai.net.	4	IN	A	23.62.96.144
a2047.w7.akamai.net.	4	IN	A	23.62.96.193
a2047.w7.akamai.net.	4	IN	A	23.62.96.162
a2047.w7.akamai.net.	4	IN	A	23.62.96.185
a2047.w7.akamai.net.	4	IN	A	23.62.96.154
a2047.w7.akamai.net.	4	IN	A	23.62.96.169
a2047.w7.akamai.net.	4	IN	A	23.62.96.152
a2047.w7.akamai.net.	4	IN	A	23.62.96.186

DNS Redirection

- **Web client's request redirected to 'close' by server**
 - **Client gets web site's DNS CNAME entry with domain name in CDN network**
 - **Hierarchy of CDN's DNS servers direct client to 2 nearby servers**



DNS Redirection

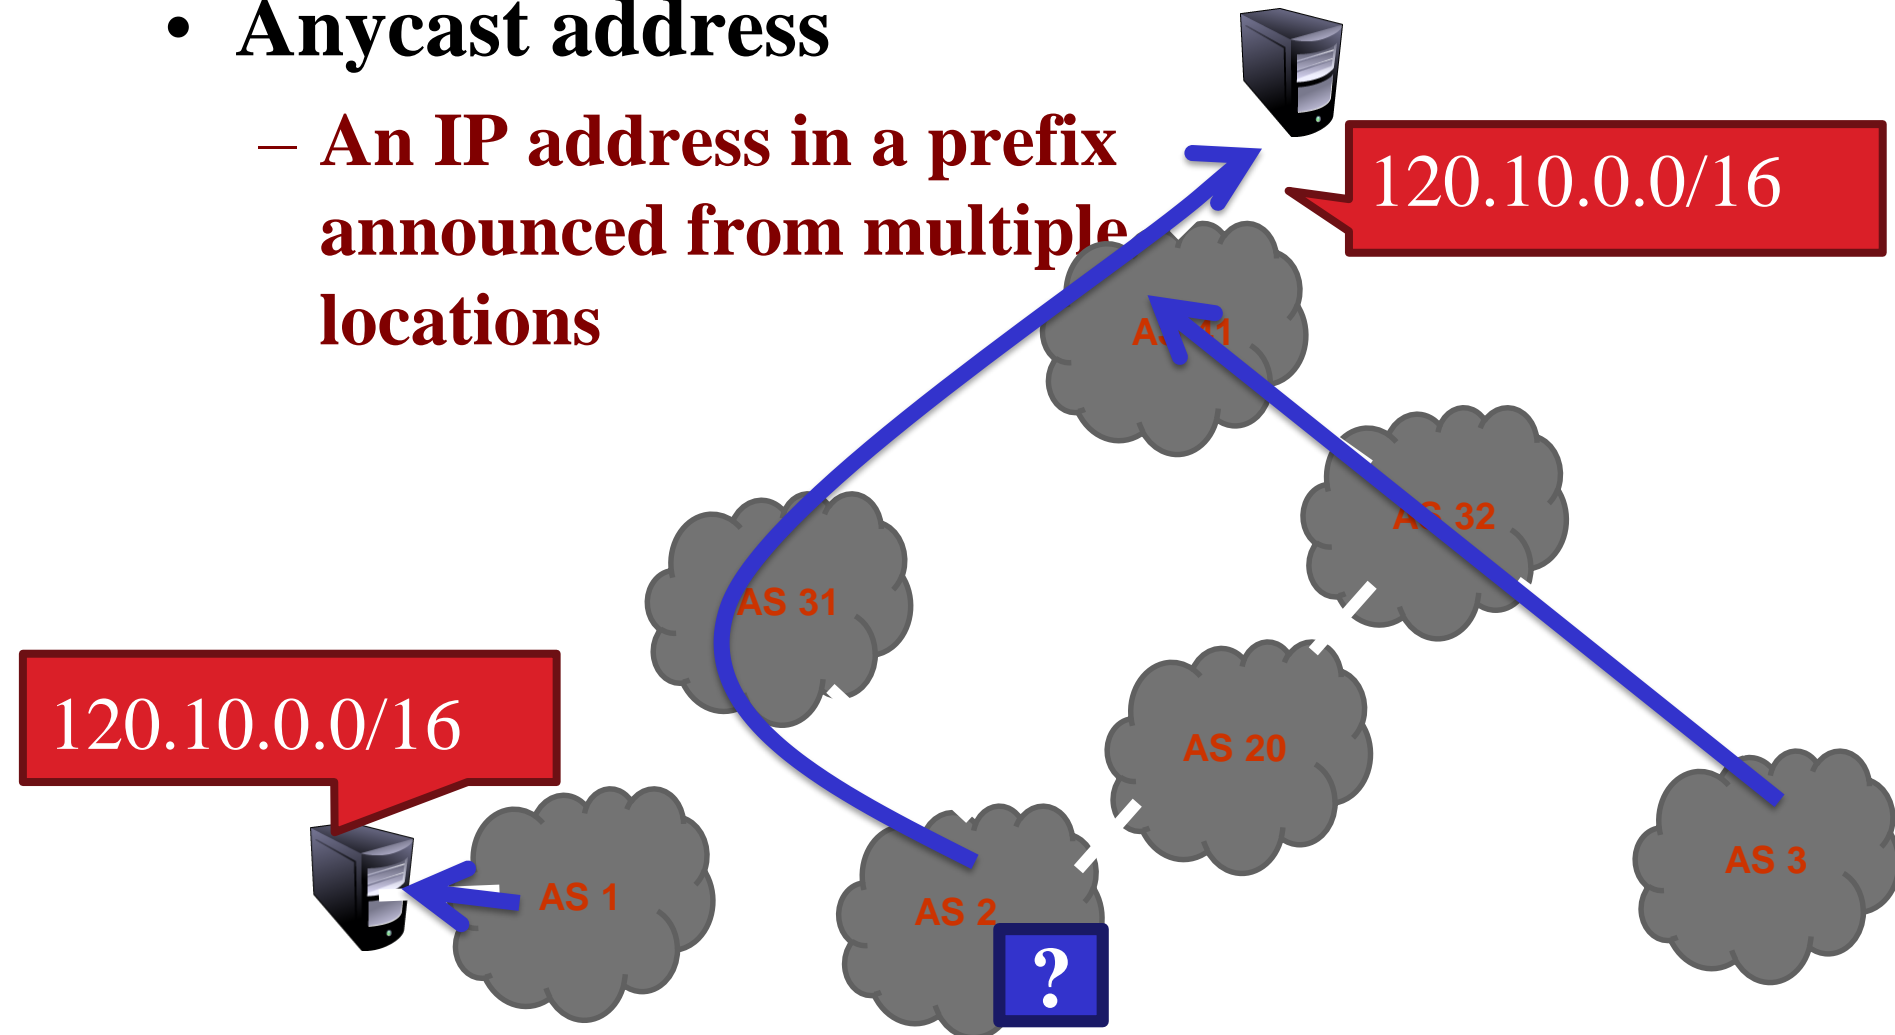


DNS Redirection Considerations

- **Advantages**
 - **Uses existing, scalable DNS infrastructure**
 - **URLs can stay essentially the same**
- **Limitations**
 - **DNS servers see only the DNS server IP**
 - Assumes that client and DNS server are close. Is this accurate?
 - **Content owner must give up control**
 - **Unicast addresses can limit reliability**

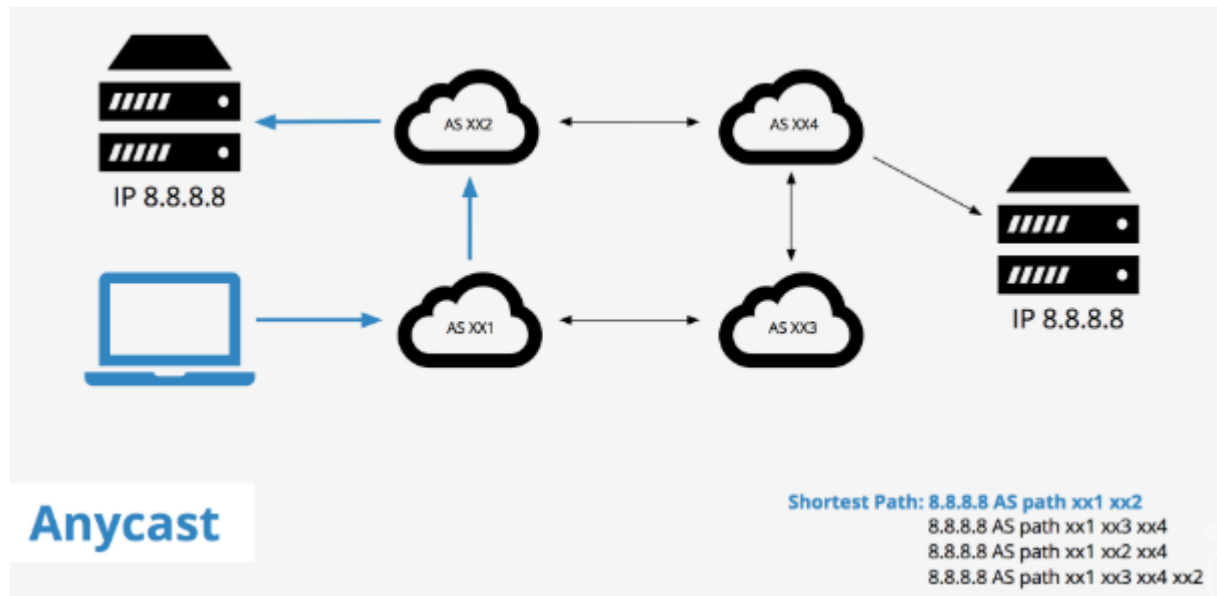
CDN Using Anycast

- Anycast address
 - An IP address in a prefix announced from multiple locations



CDN Using Anycast

- **Anycast IPv4 address /32**
- **Anycast IPv6 address /128**



Anycasting Considerations


- **Why do anycast?**
 - **Simplifies network management**
 - Replica servers can be in the same network domain
 - **Uses best BGP path**
- **Disadvantages**
 - **BGP path may not be optimal**
 - **Stateful services can be complicated**

CDN using URL Rewriting



- **URL rewriting**
 - **Origin server redirects clients to different surrogate servers by rewriting the page's URL links**

Optimizing performance


- Where to cache content?
 - **Popularity of Web objects is Zipf-like**
 - a few elements that score *very* high (the left tail in the diagrams)
 - a medium number of elements with middle-of-the-road scores (the middle part of the diagram)
 - a huge number of elements that score very low (the right tail in the diagram)
 - **Small number of sites cover large fraction of requests**
- Given this observation, how should cache replacement work?



Cache Replacement Policies (I)



Insert time: 8:00 am 7:48am 9:05am 7:10am 7:30 am 10:10am 8:45am



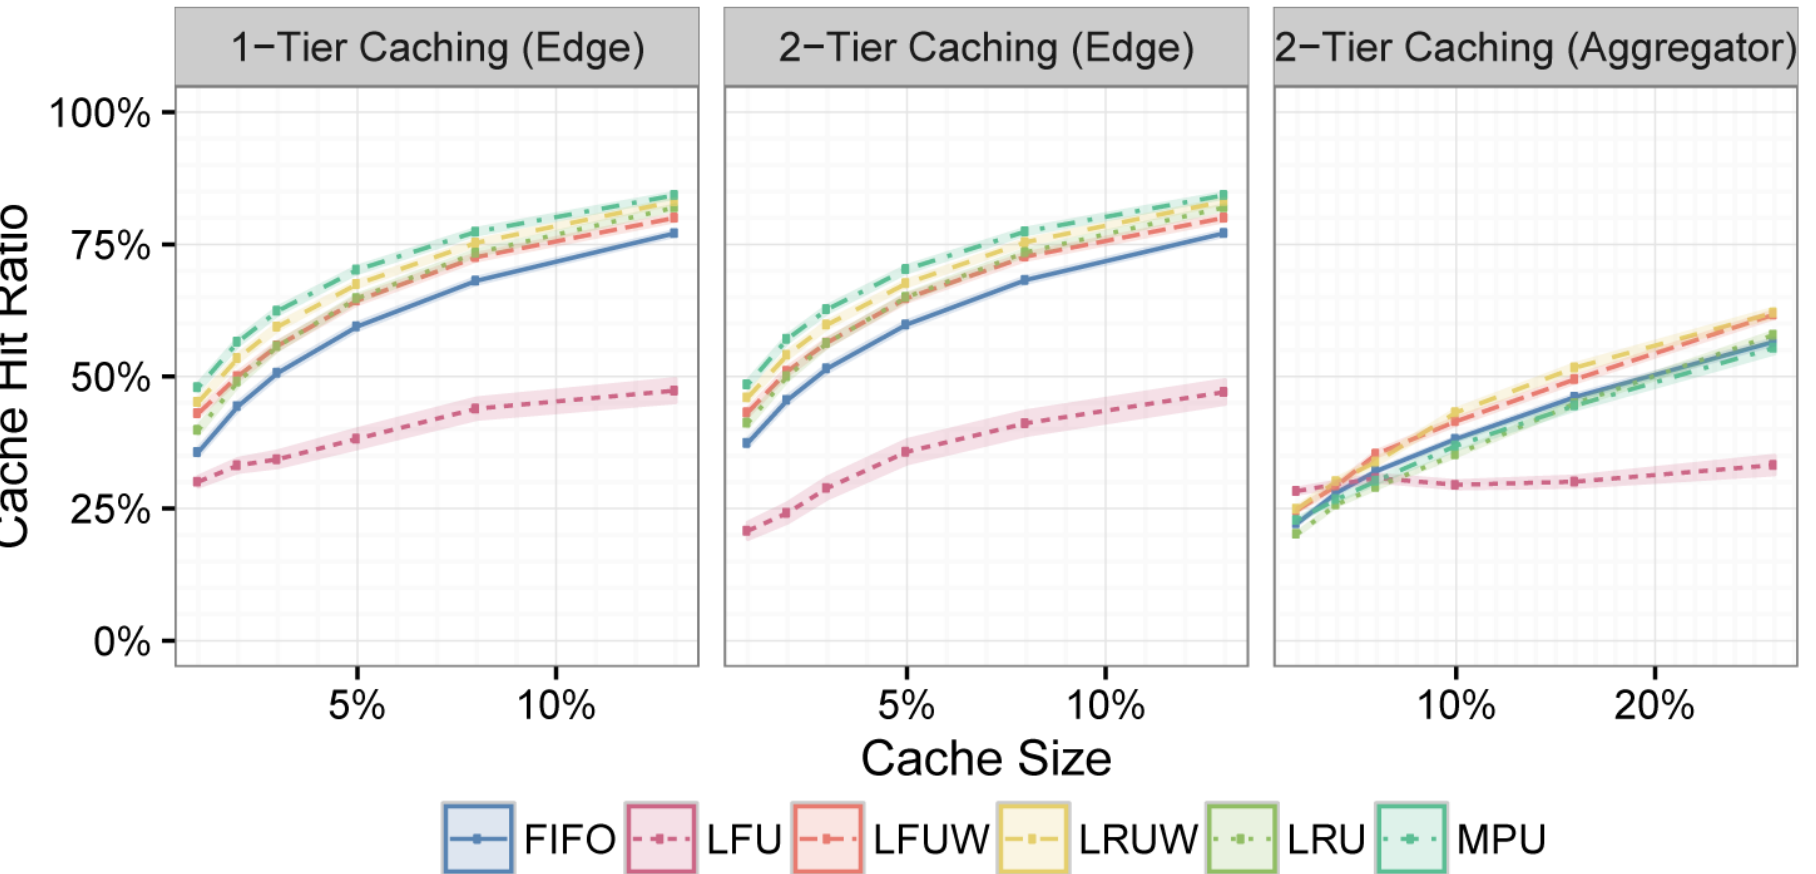
last used: 7:25am 8:12am 9:22am 6:50am 8:20am 10:02am 9:50am

Cache Replacement Policies


(II)

- **LFU: Least Frequently Used**
- **MPU: Most Probably Used**
- **LFU and LRU weighted (give a weight to each page)**


Cache Hit Ratio vs Cache Size



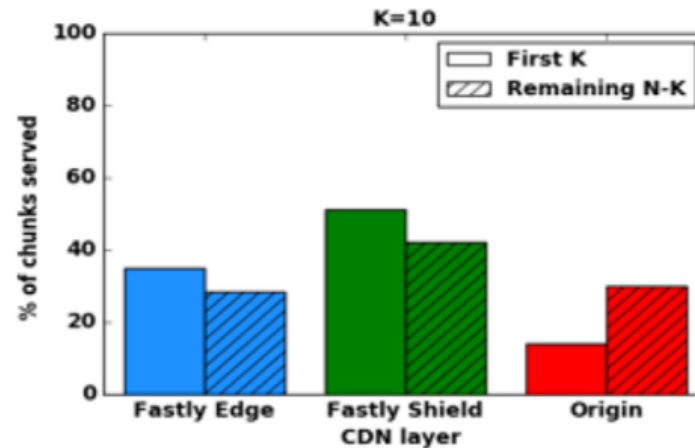
Where are objects coming from (even from chunk to chunk)?



Twitch: CloudFront/Fastly.



Vimeo: Akamai.



Vimeo: Fastly.


Peer-to-peer networks

Peer to peer networks

- **Exploits diverse connectivity between participants in a network**
- **Exploits the cumulative bandwidth of network participants**
- **Typically used for connecting nodes via large ad-hoc connections**
 - **Sharing content files containing audio, video, data**
 - **Even real-time data, such as telephony traffic, is also passed using P2P technology**
- **Pure peer-to-peer network**
 - **There is no notion of clients or servers**
 - **Equal peer nodes that simultaneously work as both "clients" and "servers" to the other nodes on the network**


The Web Model

- **Contact a server and download a web page**
- **Server has all the resources and capabilities**



The P2P Model

- A peer's resources is similar to the resources of the other participants
- P2P – peers communicating directly with other peers and sharing resources
- P2P services
 - **Distributed Computing**
 - **File Sharing**
 - **Collaboration**



Advantages

- **Clients provide resources, including bandwidth, storage space, and computing power**
- **As nodes arrive and demand on the system increases, the total capacity of the system also increases**
- **Distributed nature also increases robustness in case of failures by replicating data over multiple peers**
 - **Enable peers to find the data without relying on a centralized index server**

P2P applications

- **File sharing**
 - **Using application layer protocols**
 - DirectConnect (centralized), Gnutella (flooding), BitTorrent (hybrid)
- **VoIP**
 - **Using application layer protocols**
 - SIP
- **Streaming media**
- **Instant messaging**
- **Software publication and distribution**
- **Media publication and distribution**
 - **radio, video**

Challenges

- Peer discovery and group management
- Data **location, searching and placement**
 - **Search and routing**
- Reliable and efficient file delivery
- Security/privacy/anonymity/trust

P2P types

- **Pure P2P** refers to an environment where all the participating nodes are peers
 - No central system controls, coordinates, or facilitates the exchanges among peers
- **Hybrid P2P** refers to an environment where there are servers which enable peers to interact with each other
 - The degree of central system involvement varies with the application
 - Different peers may have different functions (simple nodes, routers, rendezvous)

No method is better than the other

- each has its advantages and its drawbacks, each is the right choice for some applications

Simple Peers

- Single end user, allowing that user to provide services from his device and consuming services provided by other peers on the network.
 - Will usually be located behind a firewall, separated from the network at large; peers outside the firewall will probably not be capable of directly communicating with the simple peer located inside the firewall.
 - Because of their limited network accessibility, simple peers have the least amount of responsibility in any P2P network.
- They are not responsible for handling communication on behalf of other peers or serving third-party information for consumption by other peers.

Rendezvous Peers

- Gathering or meeting place
 - Provides peers with a network location to use to discover other peers and peer resources.
- Peers issue discovery queries to a rendezvous peer, and the rendezvous provides information on the peers it is aware of on the network.
- May cache information on peers for future use or by forwarding discovery requests to other rendezvous peers.
 - Improve responsiveness, reduce network traffic, and provide better service to simple peers.
- Usually outside a private internal network's firewall. A rendezvous could exist behind the firewall, but it would need to be capable of traversing the firewall using either a protocol authorized by the firewall or a router peer outside the firewall.

Router (Relay) Peers


- A router peer provides a mechanism for peers to communicate with other peers separated from the network by firewall or Network Address Translation (NAT) equipment.
- Peers outside the firewall to communicate with a peer behind the firewall, and vice versa.
- A relay is not necessarily a rendez-vous peer
 - Relay is on the data stream
 - Rendez-vous is always on the discovery path (and maybe in the data stream).

Structured vs Unstructured


- Unstructured P2P networks
 - Formed when the overlay links are established arbitrarily.
 - If a peer wants to find a desired piece of data in the network, the query has to be flooded through the network to find as many peers as possible that share the data
 - The queries may not always be resolved
 - If a peer is looking for rare data shared by only a few other peers, then it is highly unlikely that search will be successful
 - Flooding causes a high amount of signalling traffic in the network
 - Gnutella and FastTrack/KaZaa, BitTorrent
- Structured P2P networks
 - Globally consistent protocol (logic) to ensure that any node can efficiently route a search to some peer that has the desired file, even if the file is extremely rare
 - The most common type of structured P2P network is the Distributed Hash Table (DHT)
 - A variant of consistent hashing is used to assign ownership of each file to a particular peer
 - Chord, Pastry, Tapestry, CAN, Tulip, Kademlia, BitTorrent (trackerless)

Fully Decentralized Information Systems

- P2P file sharing
 - Global scale application
- Example: Gnutella
 - 40.000 nodes, 3 Mio files (August 2000)
- Strengths
 - Good response time, scalable
 - No infrastructure, no administration
 - No single point of failure
- Weaknesses
 - High network traffic
 - No structured search
 - Free-riding



Gnutella: Meeting Peers (Ping/Pong)



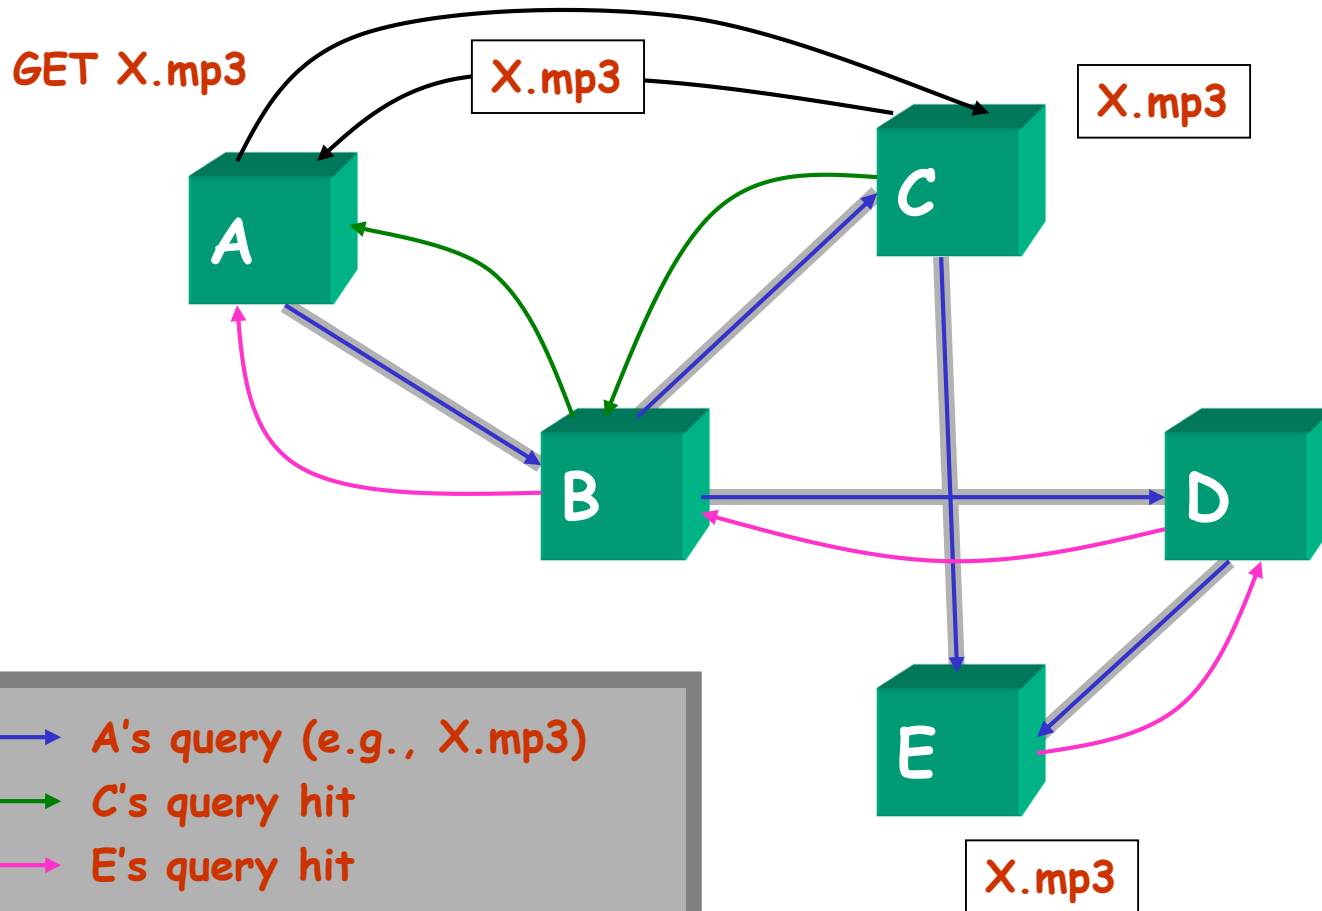
- A's ping
- B's pong
- C's pong
- D's pong
- E's pong

Gnutella: Protocol Message Types

Type	Description	Contained Information
Ping	Announce availability and probe for other servents	None
Pong	Response to a ping	IP address and port# of responding servent; number and total kb of files shared
Query	Search request	Minimum network bandwidth of responding servent; search criteria
QueryHit	Returned by servents that have the requested file	IP address, port# and network bandwidth of responding servent; number of results and result set
Push	File download requests for servents behind a firewall	Servent identifier; index of requested file; IP address and port to send file to


Gnutella: Searching

(Query/QueryHit/GET)




Searching in Gnutella (structureless)

- Queries are flooded to neighbours, have a TTL, and are forwarded only once
- Query may obtain several responses indicating which peers provides the requested file. Among those it selects one, and directly contacts it in order to download the file.
 - Can we search using fewer packets?




Improvements of Message Flooding

- **Expanding Ring**
 - start search with small TTL (e.g. $TTL = 1$)
 - if no success iteratively increase TTL (e.g. $TTL = TTL + 2$)
- **k-Random Walkers**
 - forward query to one randomly chosen neighbor only, with large TTL
 - start k random walkers
 - random walker periodically checks with requester whether to continue




Hybrid Gnutella: “Ultrapeers”

- **Ultrapeers can be installed (KaZaA) or self-promoted (Gnutella v.2)**



Real Gnutella Network

Oct 2003 Crawl of public gnutella (v.2)



**Ultrapeer
nodes**

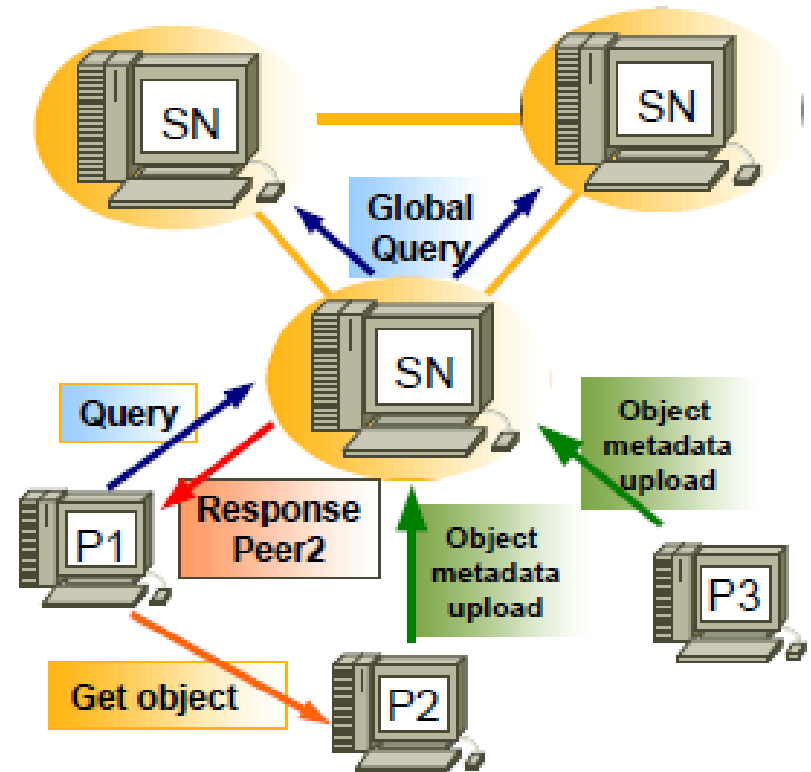
Leaf nodes

- **>100 Files**
- **0 Files**
- **0-100 Files**

- **Popular open-source file-sharing network**
 - ~450,000 users as of 2003
 - ~2,000,000 today
- **Ultrapeer-based Topology**
 - **Queries flooded among ultrapeers**
 - **Leaf nodes shielded from query traffic**
 - **Based on multiple crawlers**


FastTrack/KaZaA

- It is an extension of the Gnutella protocol which adds super-nodes to improve scalability (~gnutella v.2)
 - A peer application hosted by a powerful machine with a fast network connection become automatically a super-node, effectively acting as a temporary indexing server for other slower peers
 - Communicate between each others in order to satisfy search requests
- Network architecture: Hybrid Unstructured.
- Algorithm: Flooded Requests Model (FRM)



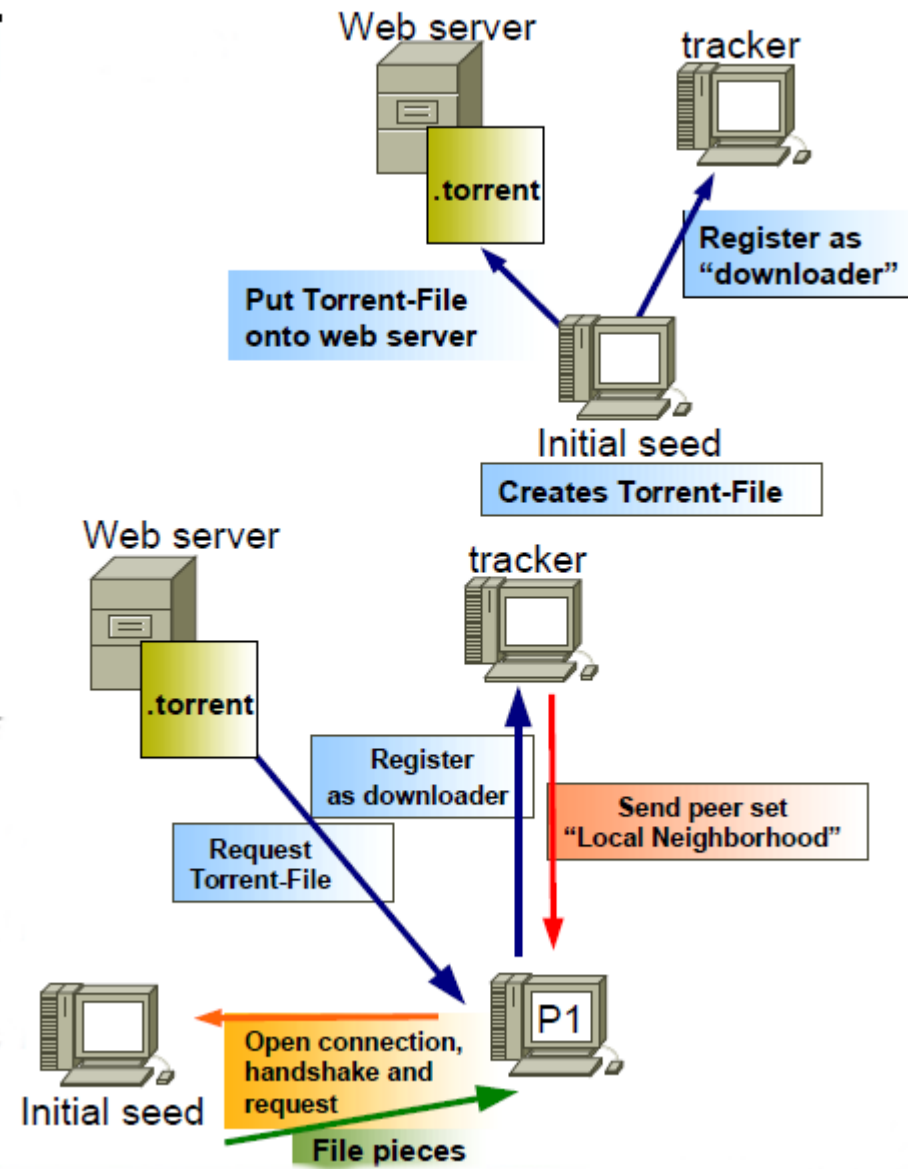
OpenNAP/Napster

- Files (music) are on the client machine
- Servers provide search (rendezvous) and initiate direct transfers between clientes
- OpenNAP is an extension to other types and linking servers.
- Network architecture: Hybrid Unstructured.
- Algorithm: Centralized Directory Model (CDM)




BitTorrent

- Trackers: keeps track of the number of seeds/peers; responsible for helping downloaders find each other, using a simple protocol on top of HTTP.
- Downloader sends status info to trackers, which reply with lists of contact information for peers which are downloading the same file.
- Web servers don't have information about content location
 - Only store metadata files describing the objects (length, name, etc.) and associating to each of them the URL of a tracker
- Network architecture: Hybrid unstructured
- Algorithm: Centralized Directory Model (CDM)
- "trackerless" torrents through a system called the "distributed database", through DHT (Distributed Hash Tables)

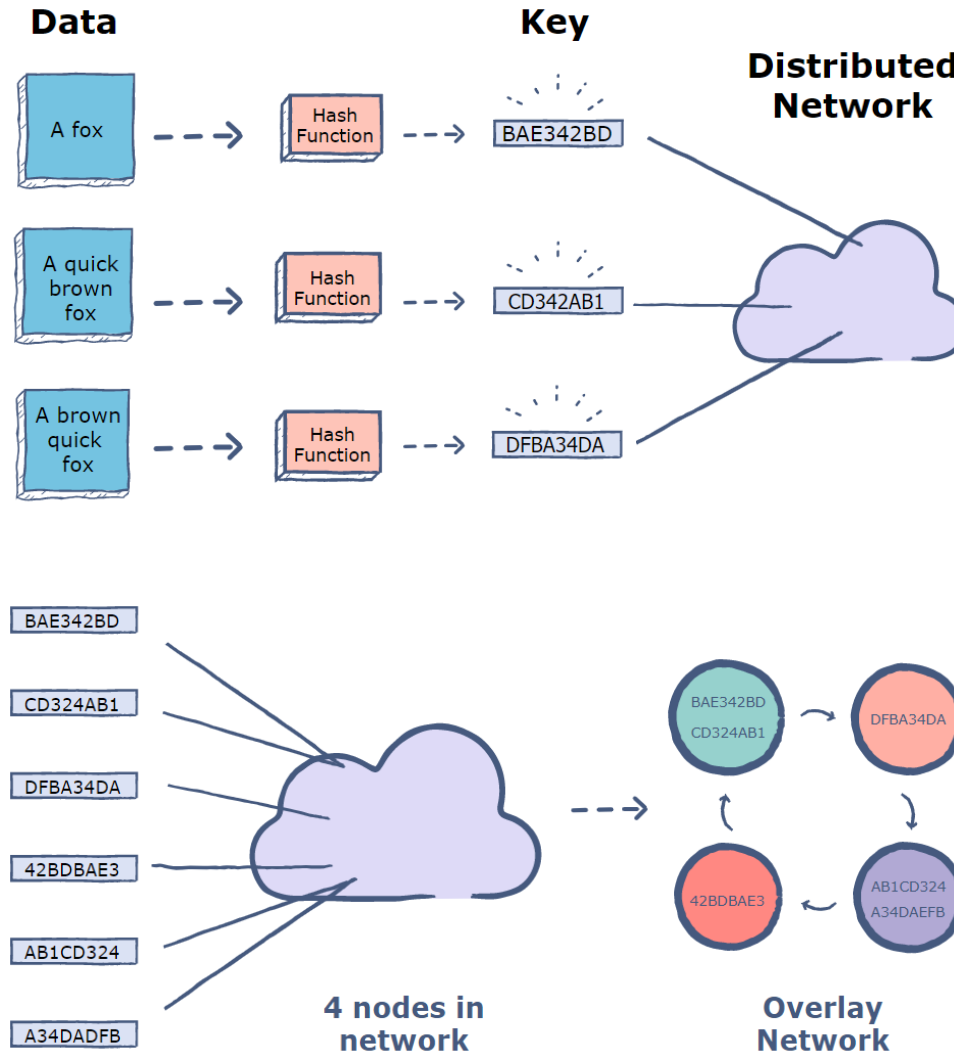


Searching in DHTs (structured)

- Need to know the exact filename
 - **Keys (filenames) map to node-ids**
 - Change in file name → search at different nodes
 - No wildcard matching: cannot ask for file “pix*”




DHT

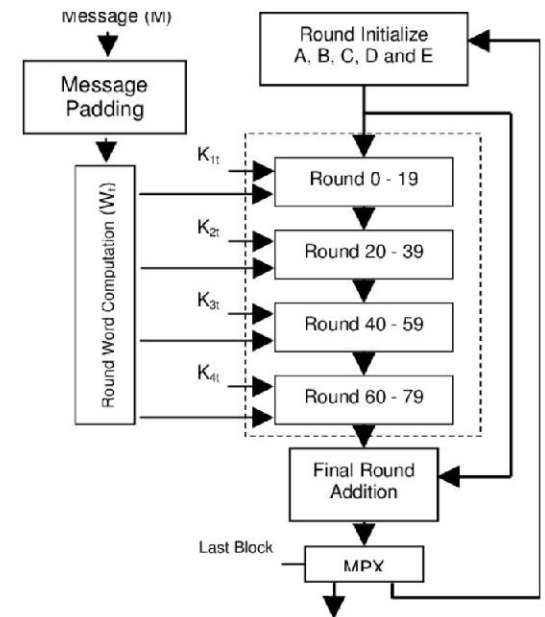


CHORD: DHT Algorithm

- All files/data items in the network will have an identifier, which will be hashed to give a key for that particular resource
- If a node needs a file/data, it will hash its name and then send a request using this key.
- All n nodes also use the function to hash their IP address, and conceptually, the nodes will form a ring in ascending order of their hashed IP



SHA-1




CHORD: DHT Algorithm

- The successor node of a key k is the first node whose ID equals to k or follows k in the identifier circle, denoted by $\text{successor}(k)$
- Every key is assigned to its successor node, so looking up a key k is to query $\text{successor}(k)$.
- When a node wants to share a file or some data
 - Hashes the identifier to generate a key k , and sends its IP and the file identifier to $\text{successor}(k)$
 - These are then stored at this node
 - All resources are indexed in a large DHT across all participating nodes
 - If there are two or more nodes that hold a given file or resource, the keys will be stored at the same node in the DHT, giving the requesting node a choice

DHT: Store Information

- Hashing of search keys AND peer addresses on binary keys of length m
 - e.g. $m=8$, $\text{key}(\text{'jingle-bells.mp3'})=17$, $\text{key}(196.178.0.1)=3$
- Data keys are stored at next larger node key



Search possibilities


1. every peer knows every other
 $O(n)$ routing table size
2. peers know successor
 $O(n)$ search cost

DHT: Store Information

- Every peer knows m peers with exponentially increasing distance

Each peer p stores a routing table

First peer with hashed identifier s_i such that $s_i = \text{successor}(p + 2^{i-1})$ for $i = 1, \dots, m$



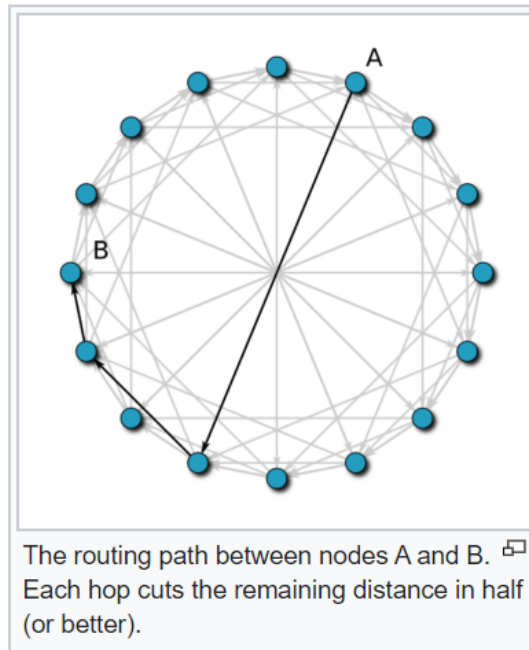
i	s_i
1	$p2$
2	$p2$
3	
4	$p3$
5	$p4$

Search

$O(\log n)$ routing table size

DHT: Search Information


- When a node wants a content
 - Hashes its identifier and sends a request to $\text{successor}(k)$
 - Reply with the IP of the node that holds the actual data
 - How does a node request information from $\text{successor}(k)$, when it doesn't know its IP, but only the key?
 - Every node holds what is known as a finger table
 - Contains a list of keys and their successor IP's
 - Each node holds the IP of an exponential sequence of nodes that follow it, i.e. entry i of node k 's finger table holds the IP of node $k + 2^i$



Search

`search(p, k)`

find in routing table largest (i, p^*) such that $p^* \in [p, k[$
 /* largest peer key smaller than the searched data key */
 if such a p^* exists then `search(p^* , k)`
 else return (`successor(p)`) // found



Search

$O(\log n)$ search cost

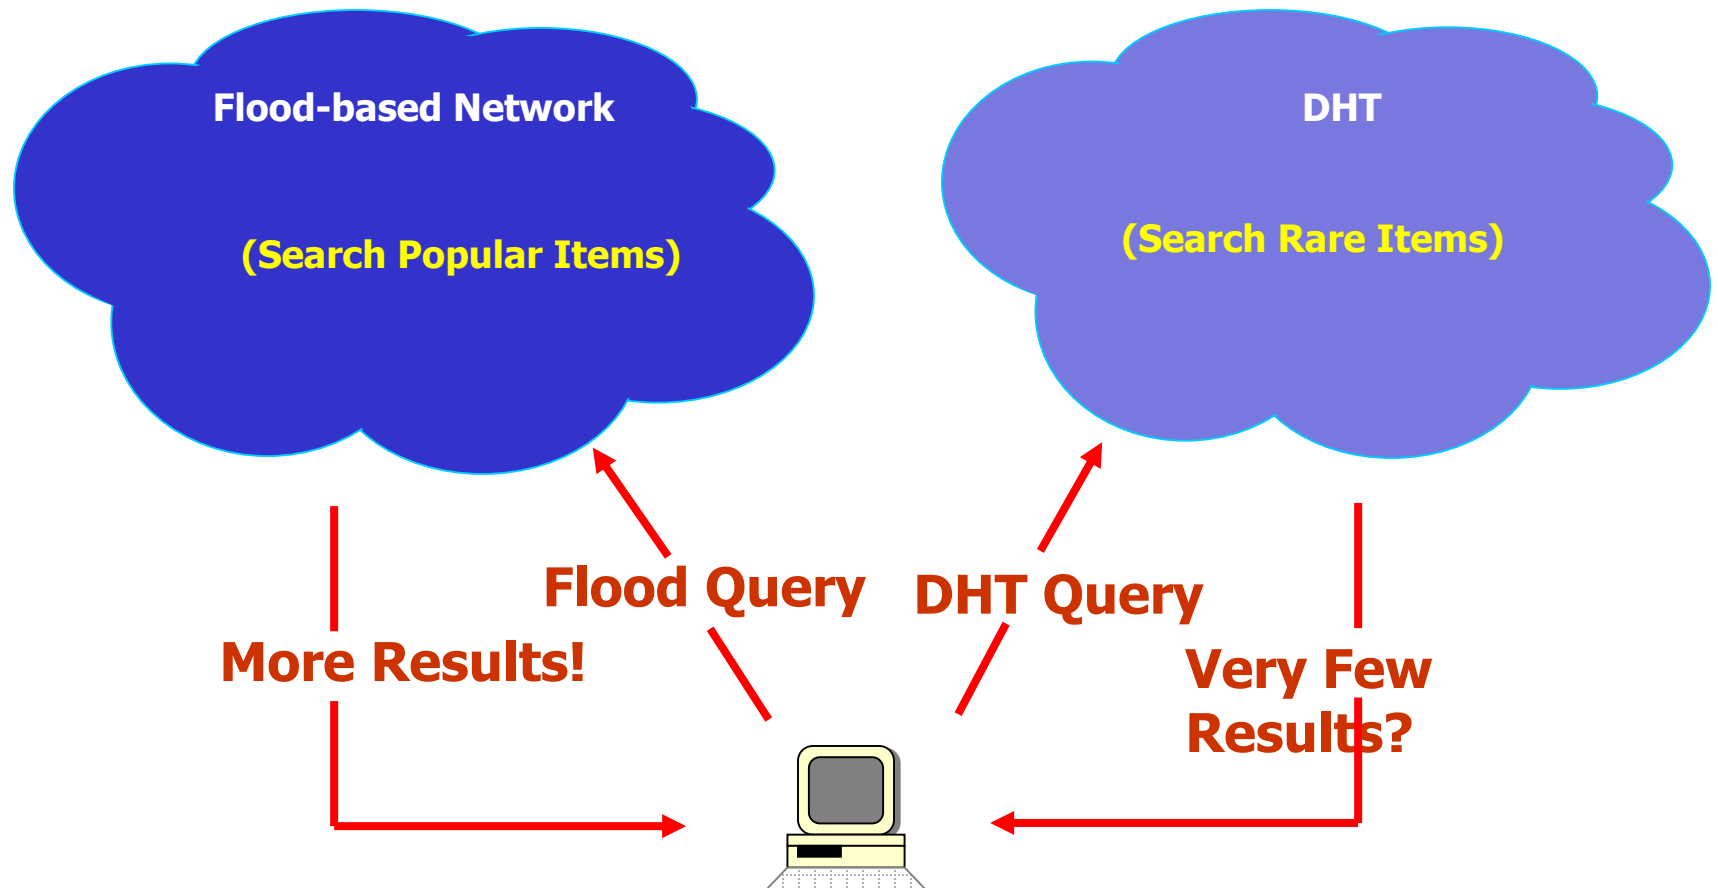
Routing Table with exp.
 increasing distance $\Rightarrow O(\log n)$
 with high probability

File Search: Flooding vs. DHTs

- **Recall**
 - **Flooding can miss files**
 - **DHTs should never**
- **Query complexity**
 - **Flooding can handle arbitrary single-site logic**
 - **DHTs can do equijoins, selections, aggregates, etc.**
 - But not so good at fancy selections like wildcards
- **Query Performance**
 - **Flooding can be slow to find things, uses lots of BW**
 - **DHTs: expensive to publish documents with lots of terms**
 - **DHTs: expensive to intersect really long term lists**
 - Even if output is really small!
- **Hybrid solution!**

Hybrid Search

Hybrid = “Best of both worlds”



Security - attacks

- **Poisoning attacks**
 - e.g. providing files whose contents are different from the description
- **Polluting attacks**
 - e.g. inserting "bad" chunks/packets into an otherwise valid file on the network
- **Freeloaders**
 - Users or software that make use of the network without contributing resources to it
- **Insertion of viruses to carried data**
 - e.g. downloaded or carried files may be infected with viruses or other malware
- **Malware in the peer-to-peer network software itself**
 - e.g. distributed software may contain spyware
- **Denial of service attacks**
 - Attacks that may make the network run very slowly or break completely
- **Filtering**
 - Network operators may attempt to prevent peer-to-peer network data from being carried
- **Identity attacks**
 - e.g. tracking down the users of the network and harassing or legally attacking them
- **Spamming**
 - e.g. sending unsolicited information across the network- not necessarily as a denial of service attack

Security

- **Most attacks can be defeated or controlled by careful design of the peer-to-peer network and through the use of encryption**
 - **However, almost any network will fail when the majority of the peers are trying to damage it**
- **Anonymity**
 - **Some peer-to-peer protocols (such as Freenet) attempt to hide the identity of network users by passing all traffic through intermediate nodes**
- **Encryption**
 - **Some peer-to-peer networks encrypt the traffic flows between peers**
 - Make it harder for an ISP to detect that peer-to-peer technology is being used (as some artificially limit bandwidth)
 - Hide the contents of the file from eavesdroppers
 - Impede efforts towards law enforcement or censorship of certain kinds of material
 - Authenticate users and prevent 'man in the middle' attacks on protocols
 - Aid in maintaining anonymity

To Conclude:

P2P and self-organization in Architectures

P2P self-organization concepts can be found at every layer, reflecting some sort of self-organization in the communication structure.

Layer	Communication type	Information discovery	Information transport
<i>(at which Layer we are considering self-organization)</i>	<i>Communication concept explored</i>	<i>How to figure where the information is</i>	<i>How is the information encapsulated for transport</i>
Networking Layer	Regular Internet (IP) protocols	Routing, DNS	TCP
Data Access Layer	Overlay Networks, P2P	Resource Location (DHT, central server)	Gnutella, FreeNet
Service Layer	Application interface	Messaging, Distributed Processing	Napster, SETI, Groove
User Layer	User Communities, Google Circles	Collaboration	eBay, Google+, Facebook,