

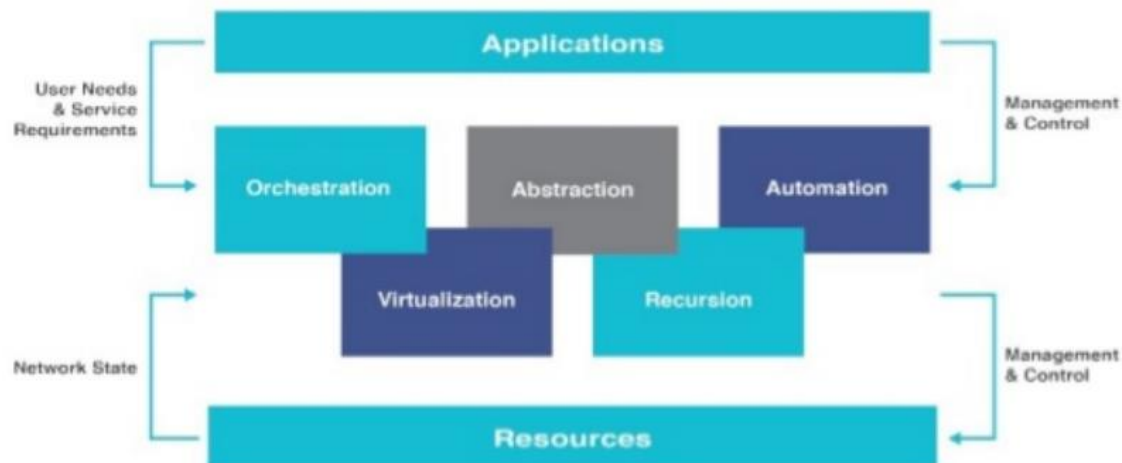
# **Software Defined Networks and OpenFlow**

**Mestrado Integrado em  
Engenharia de Computadores e  
Telemática  
2019/2020**

# Need for Programmable Networks

- Changing traffic patterns
- Rise of cloud services
- 'Big data' requires more bandwidth
- Inability to scale
- Need for open programmable networks
- Open standards-based and vendors-neutral

## The Architecture of Software-Defined Networks



# Characteristics of SDN


- Plane separation
- Simplified device
- Centralized control
- Network automation and virtualization
- Openness

# Plane Separation

- Separation of the forwarding and control planes
- Forwarding plane:
  - logic and tables for choosing how to deal with incoming packets based on characteristics such as MAC address, IP address, and VLAN ID
  - Forward, drop, consume, or replicate an incoming packet
  - Device determines the correct output port by performing a lookup in the address table in the hardware
  - Special-case packets that require processing by the control or management planes are consumed and passed to the appropriate plane

# Plane Separation

- Control plane
  - Determines how the forwarding tables and logic in the data plane should be programmed
  - Run routing or switching protocols
  - Performed in a centralized controller
  - Can be co-located in the switches



# Simplified Device and Centralized Control

- Control software is removed from the device and placed in a centralized controller
- Software-based controller manages the network using higher-level policies
- Controller provides primitive instructions to the simplified devices when appropriate
  - Allow devices to make fast decisions about how to deal with incoming packets.

# Network Automation and Virtualization

- SDN is said to be a natural evolution for the problem of network control
- Distributed state abstraction provides the network programmer with a global network view
  - Specify the necessary forwarding behaviors without any knowledge of vendor-specific hardware
  - Express the desired goals of the overall network without getting lost in the details of how the physical network will implement those goals
- Open interface on the controller to allow for automated control of the network
  - Northbound and southbound: interface is to the applications or to the devices

# Network Automation and Virtualization

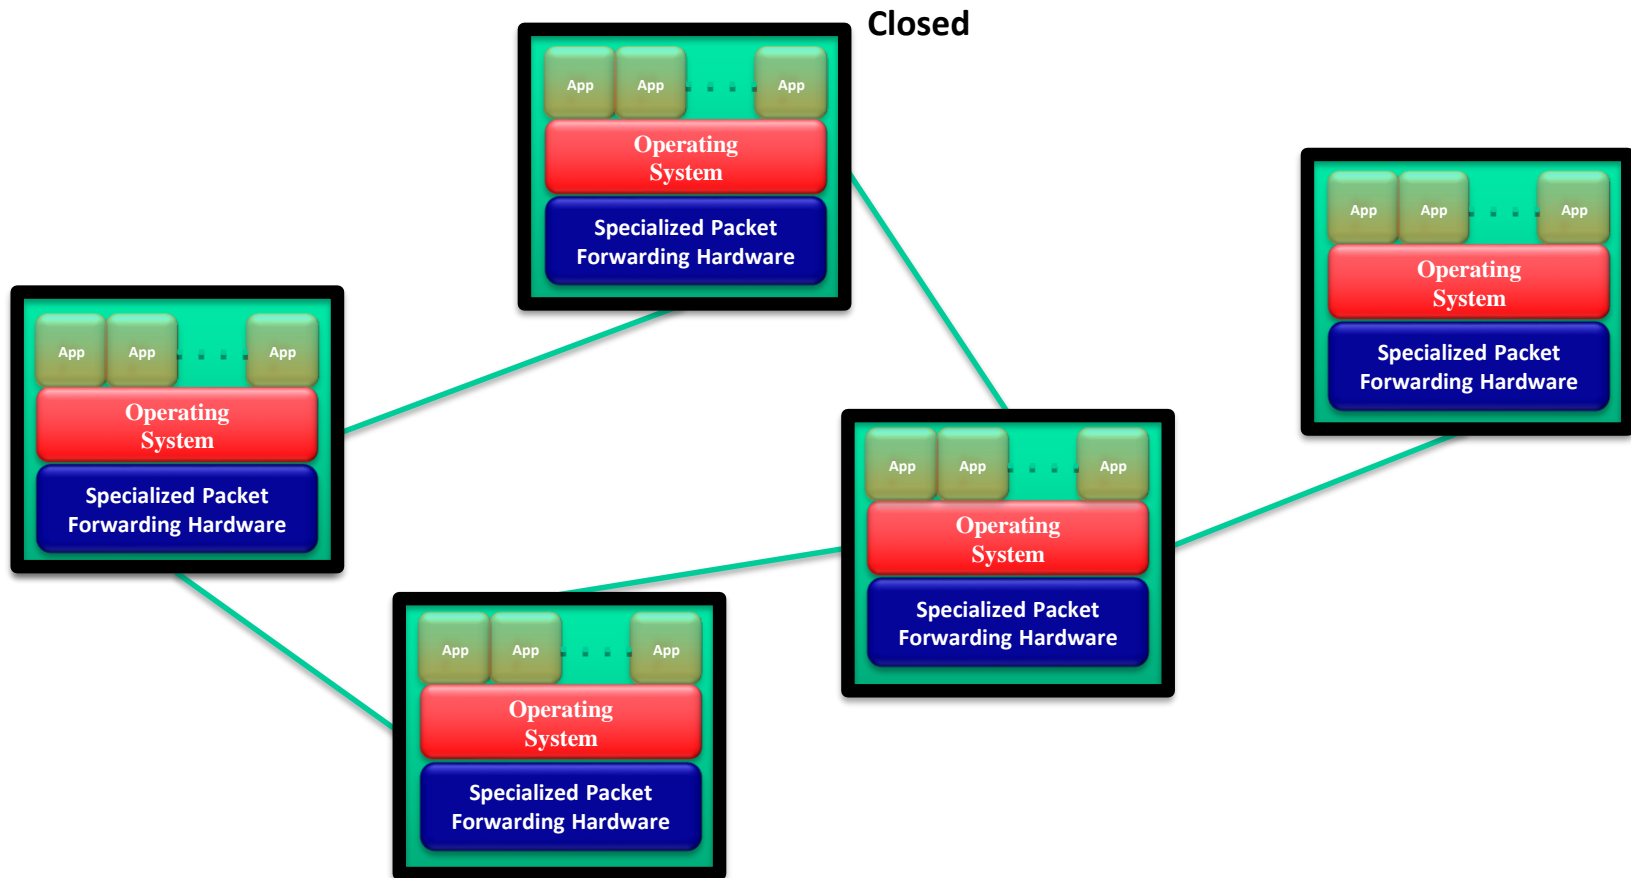
- Southbound API is the OpenFlow interface to program the network devices
- Northbound API for software applications to be plugged into the controller
  - Algorithms and protocols
  - Can quickly and dynamically make network changes as the need arises
  - Interface that allows the software above it to operate without knowledge of the individual characteristics and of the network devices themselves
  - Applications can be developed that work over a wide array of manufacturers' equipment that may differ substantially in their implementation details.
  - Ability to virtualize the network, decoupling the network service from the underlying physical network
    - Unaware if network resources are virtual or physical



# Openness

- Standard interfaces should remain standard, well documented, and not proprietary
- APIs to give software sufficient control to experiment with and control various control plane options
- Northbound and southbound interfaces
  - Easily experiment with and test new ideas
  - Resulting in better and faster technological advancement in the structure and functioning of networks
  - Open interfaces also encourages SDN-related open source projects, and permit equipment from different vendors to interoperate

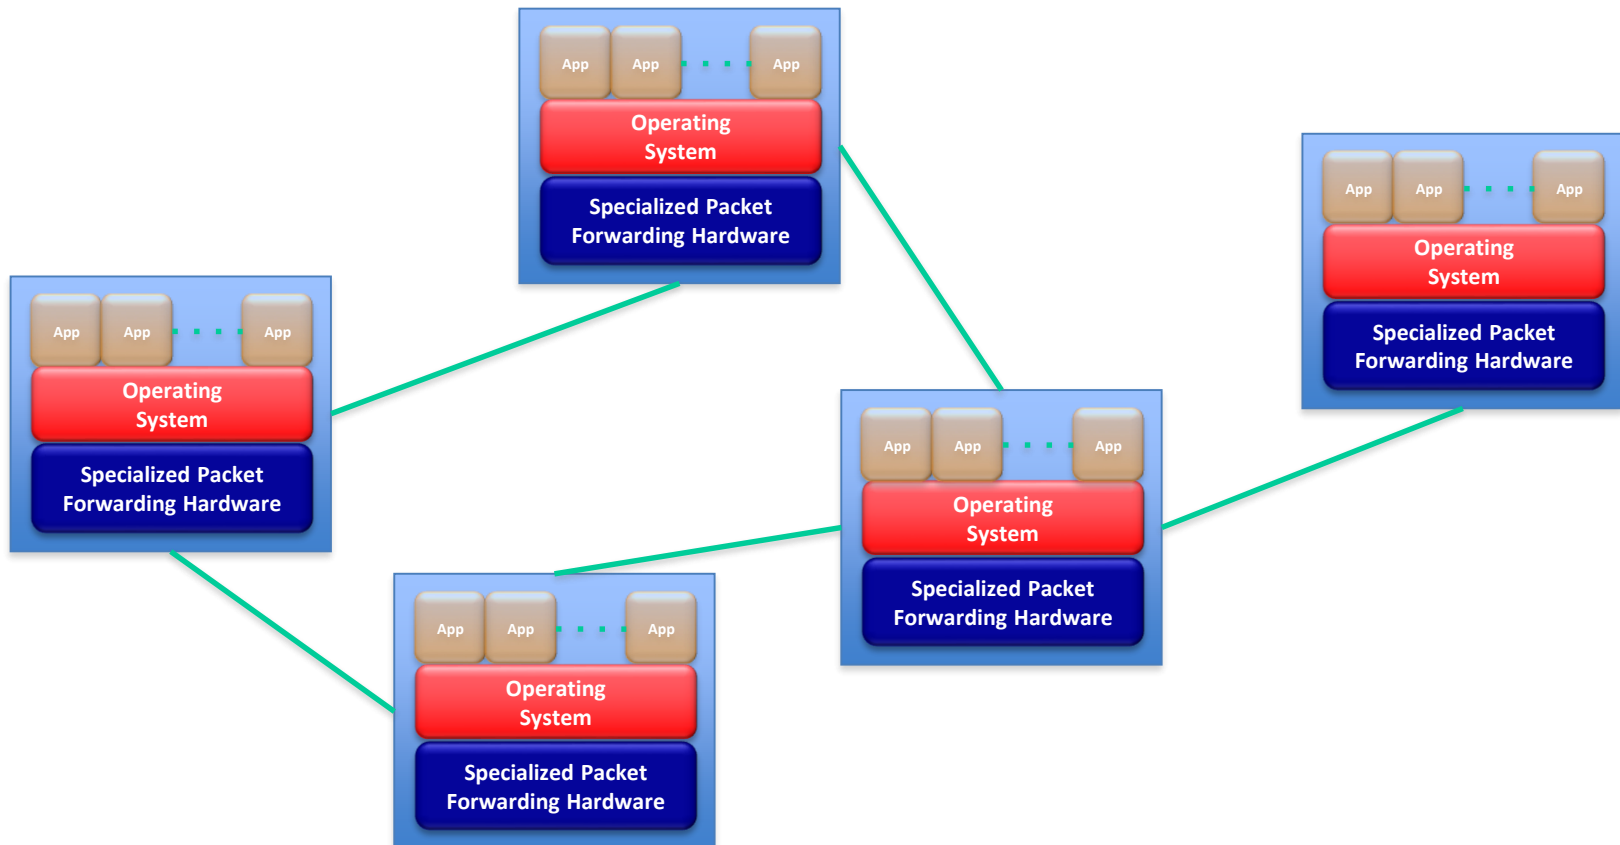
# Idea: An OS for Networks



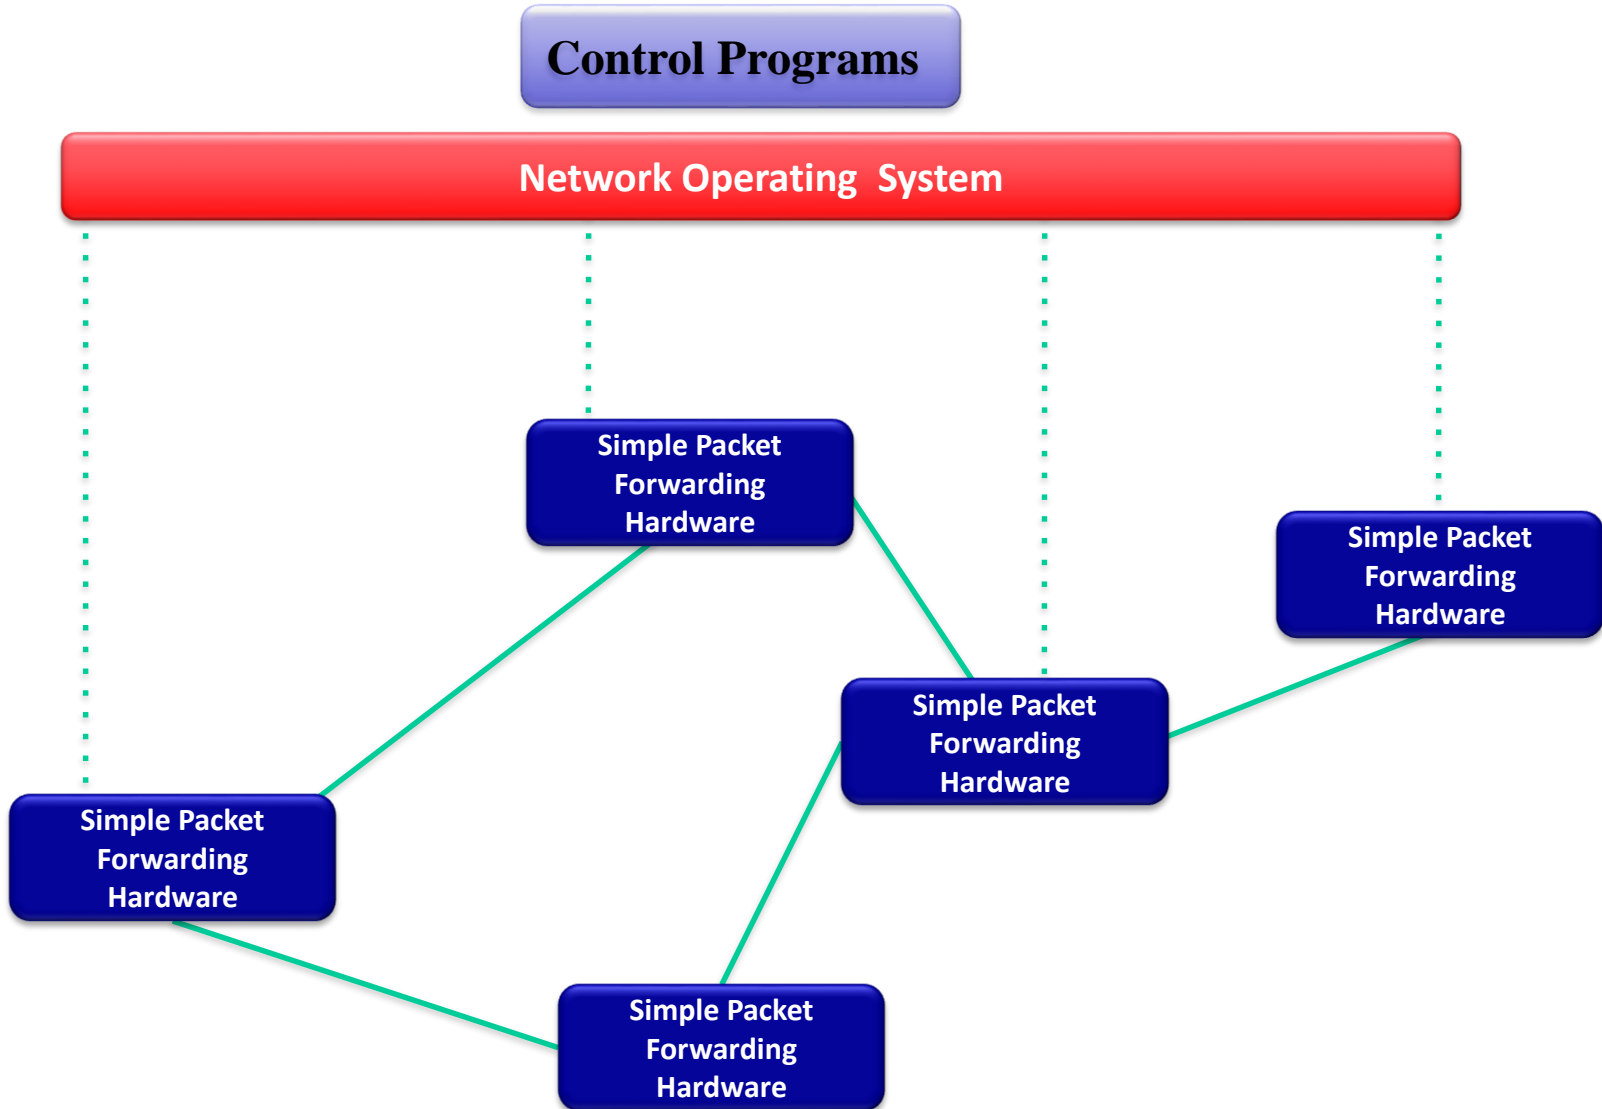
# Idea: An OS for Networks

Control Programs

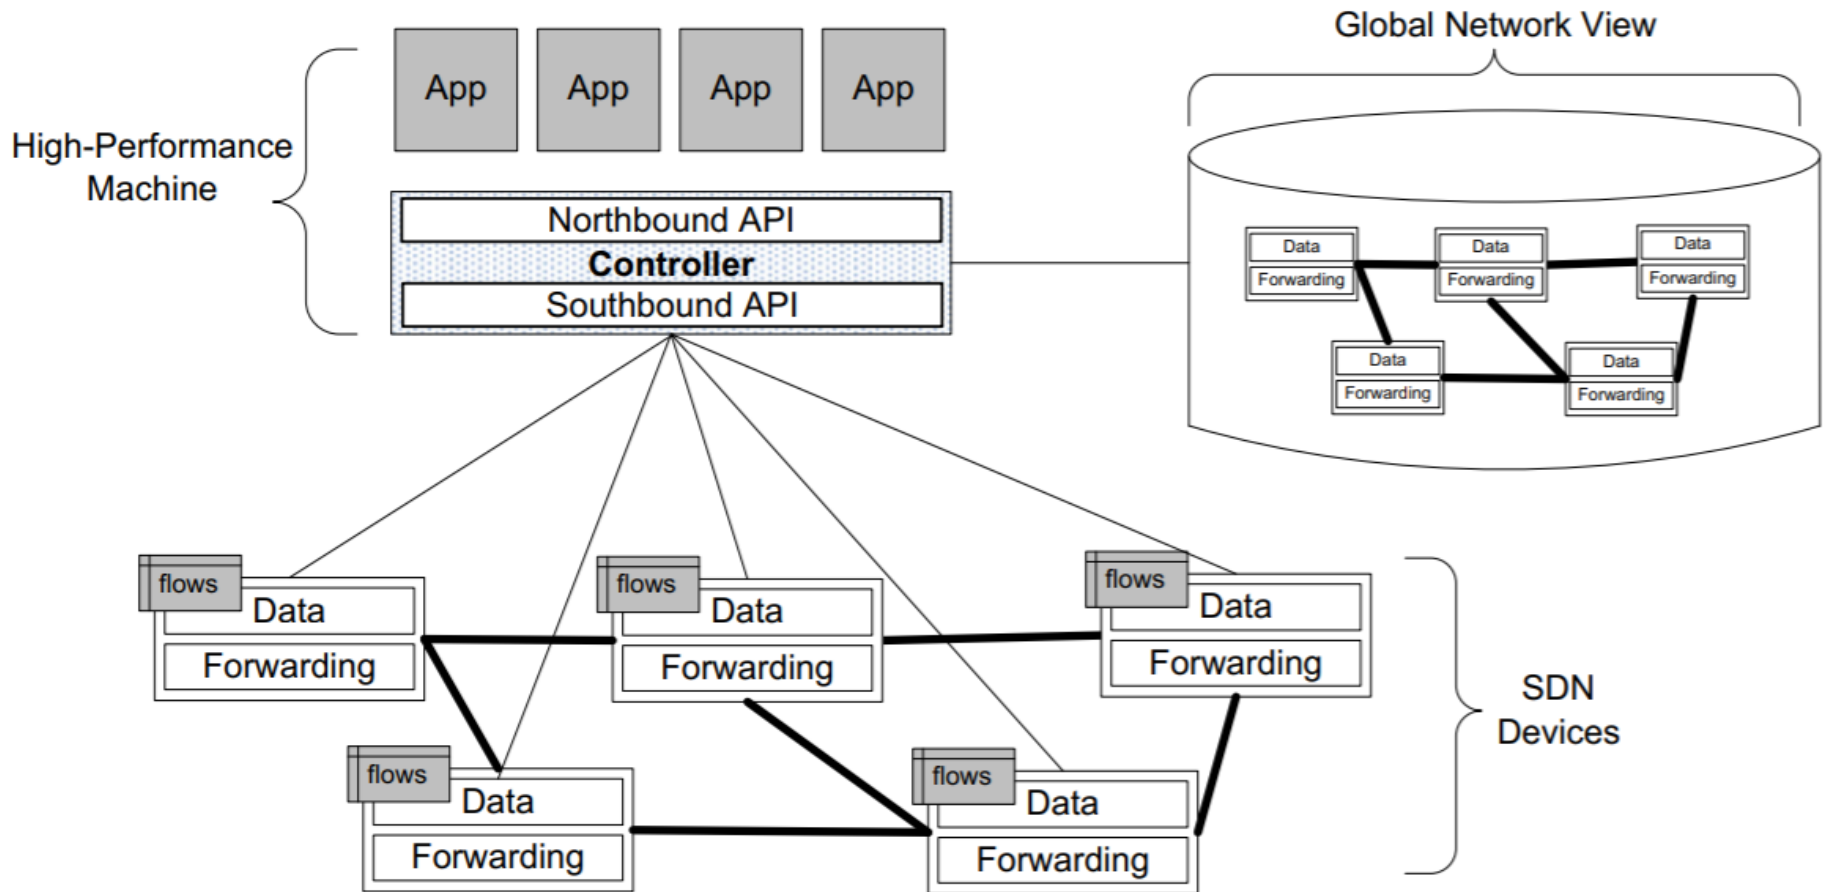
Network Operating System



# Idea: An OS for Networks



# SDN: Architecture and Operation



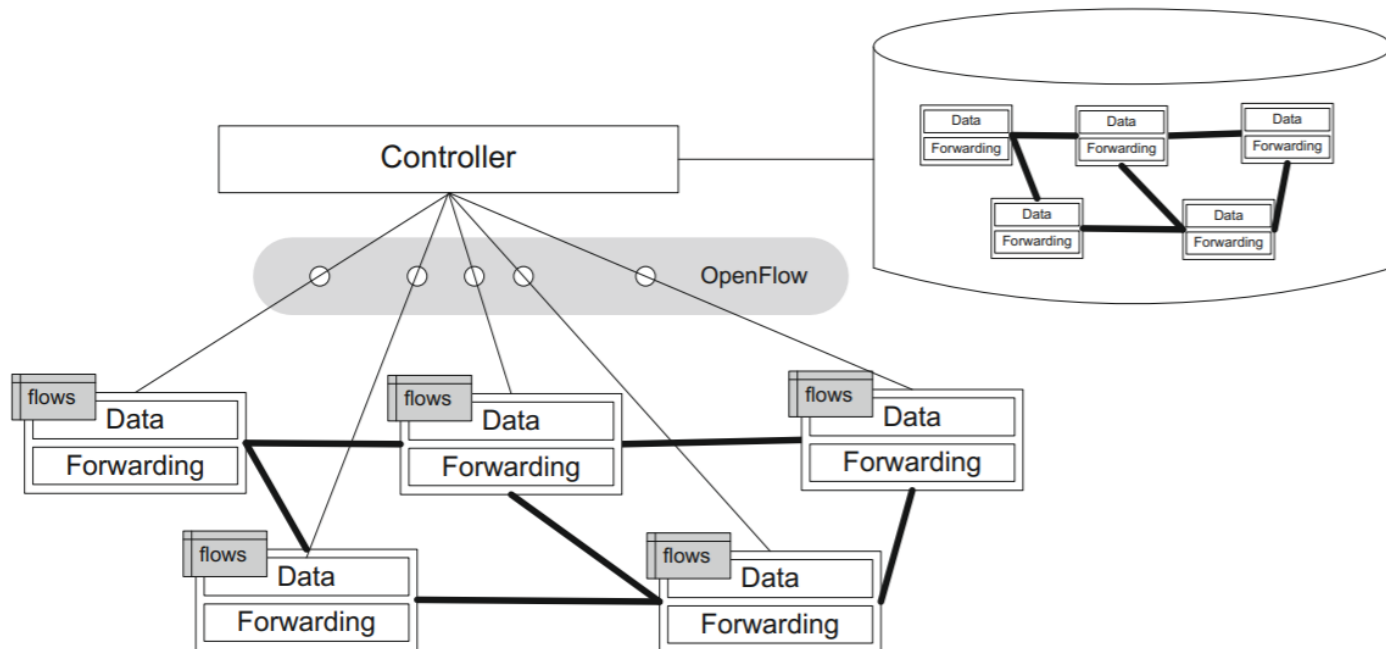
Flow describes a set of packets transferred from one network endpoint (or set of endpoints) to another endpoint (or set of endpoints).

# SDN Operation: Flows

- IP address-TCP/UDP pairs, VLAN endpoints, layer three tunnel endpoints, and input ports, among other things
- Rules describes the forwarding actions that the device should take for all packets belonging to that flow
- Flows are unidirectional, opposite direction is a separate flow
- Flow table: series of flow entries and the actions to perform when a packet matching that flow arrives at the device; if it does match any flow, it is discarded or decision does to the controller

# SDN Operation: Controller

- Allows the SDN application to define flows on devices
- Helps the application respond to packets that are forwarded to the controller by the SDN devices
- Calculates optimal forwarding solutions for the network in a deterministic, predictable manner



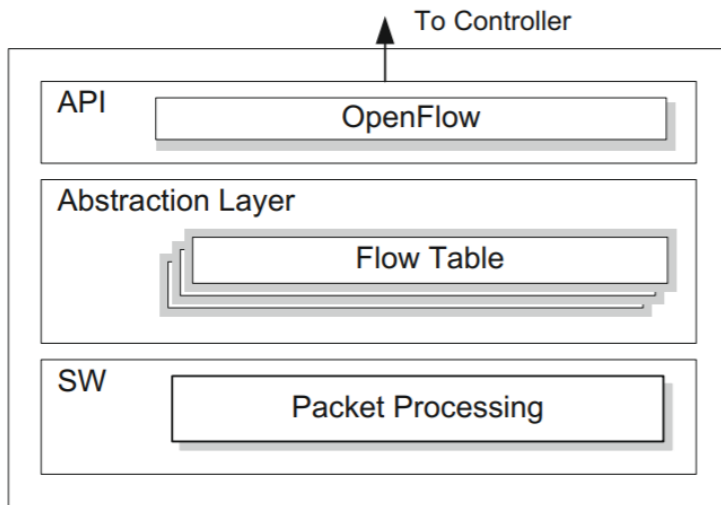
# SDN Operation: Applications

- Part of network layers two and three
- Use SDN controller to set proactive flows on the devices and to receive packets that have been forwarded to the controller
- Flows defined in response to a packet forwarded to the controller
  - SDN application will instruct the controller as to how to respond to the packet
  - If appropriate, will establish new flows on the device in order to allow that device to respond locally the next time it sees a packet belonging to that flow: reactive flows
- Software applications: forwarding, routing, overlay, multipath, and access control functions, among others
- Controller can insert flows reactively in response to other data sources such as intrusion detection systems (IDS) or traffic analyzer

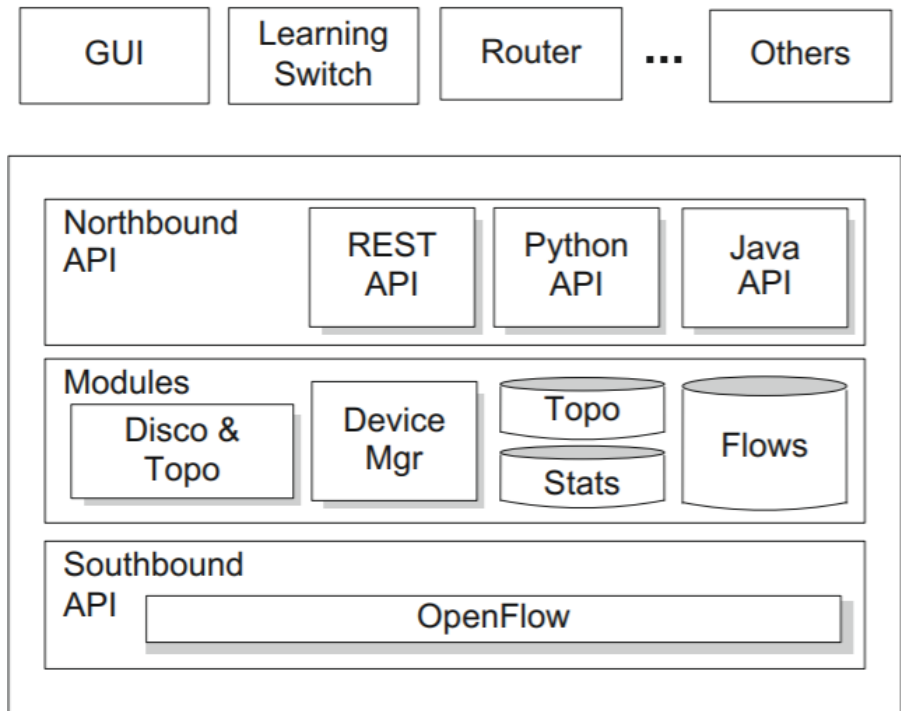


# SDN Operation: OpenFlow

- Means of communication between the controller and the device



**SDN switch**



**SDN controller**

# SDN Operation: SDN controller

- Device and topology discovery and tracking
  - Learning of the existence of switches (SDN devices) and end-user devices and tracking the connectivity between them
- Flow management, device management, and statistics tracking
  - Maintains a flow cache that mirrors the flow tables on the various switches it controls
  - Locally maintains per-flow statistics that it has gathered from its switches

# SDN Operation: Scaling the number of flows

- At the edge, flows will permit different policies to be applied to individual users and even different traffic types of the same user
  - Multiple flow entries for a single user
- Core devices, the flow definitions will be generally more coarse, with a single aggregated flow entry matching the traffic from a large number of users whose traffic is aggregated in some way, such as a tunnel, a VLAN, or a MPLS LSP
  - Policies applied deep into the network to aggregated flows

# SDN switches

- Software:
  - Open vSwitch (OVS)
  - Indigo
- Network equipment manufacturers
  - Cisco
  - HP
  - NEC
  - IBM
  - Juniper
  - Extreme
- Have added OpenFlow support to some of their legacy switches





# SDN controllers

- Software:
  - Ryu
  - Floodlight controller (Java and a RESTful API)
  - OpenDaylight controller (RESTful API)
  - NOX
  - Beacon
  - Trema
- Network equipment manufacturers
  - Cisco
  - NEC
  - IBM
  - HP

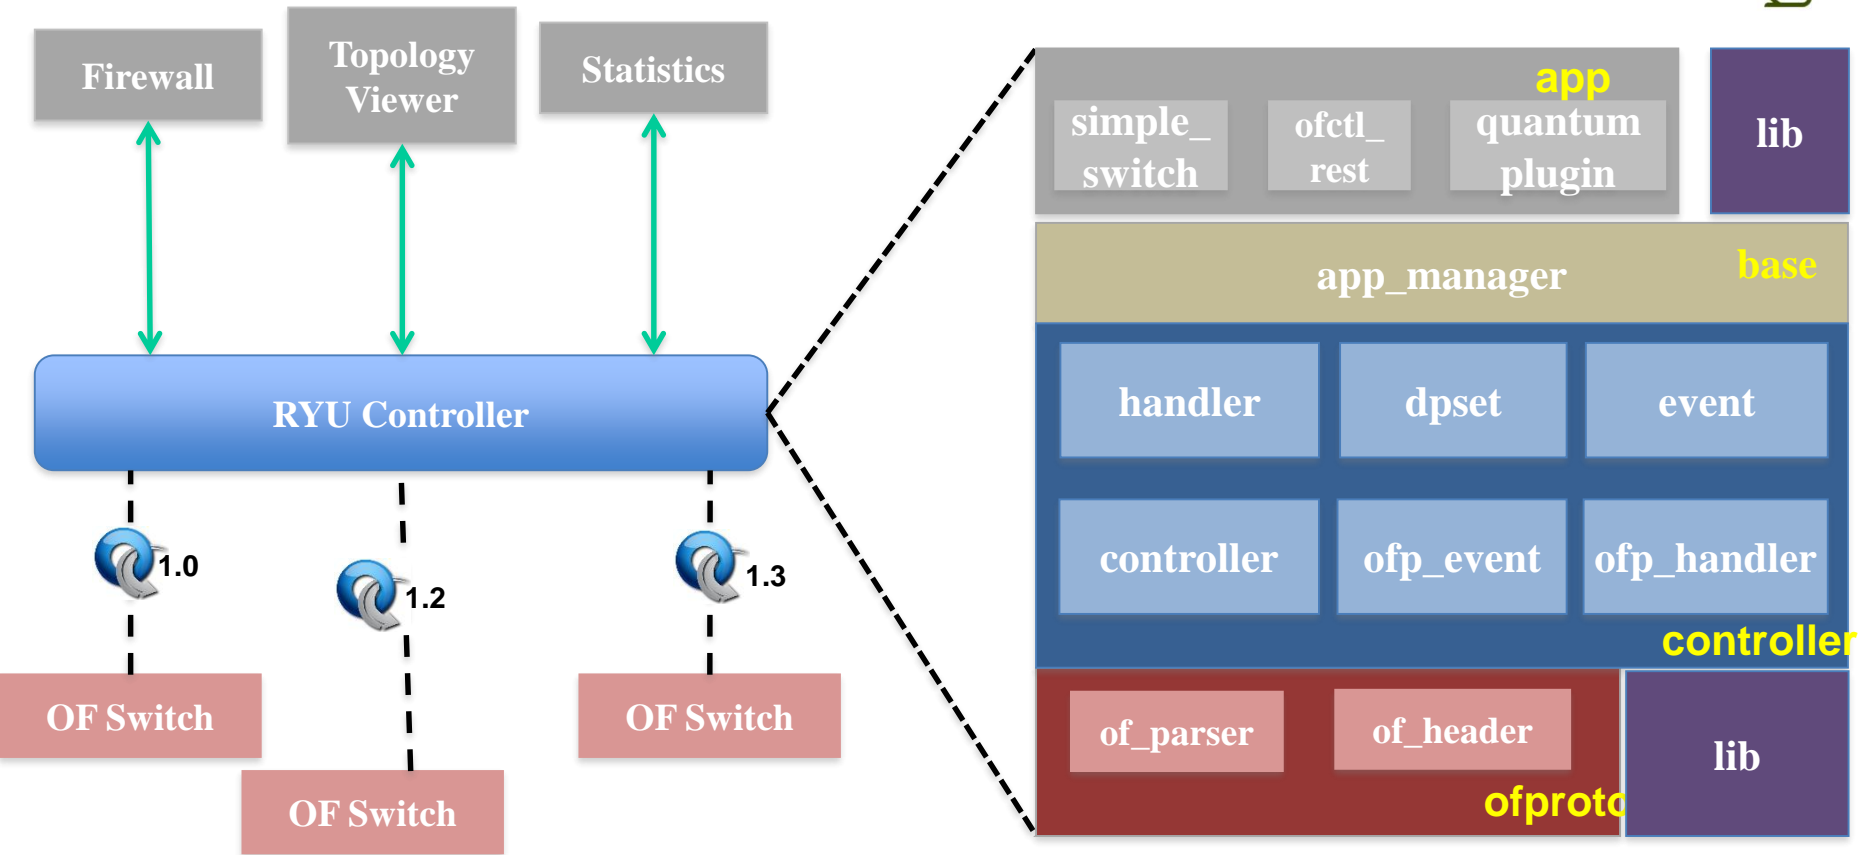
# Open-source controllers

Controller	Notes
Ryu	<ul style="list-style-type: none"><li>•Apache license</li><li>•Python</li></ul>
NOX/POX	<ul style="list-style-type: none"><li>•GPL</li><li>•C++ and Python</li></ul>
Stanford's Beacon	<ul style="list-style-type: none"><li>•BSD-like license</li><li>•Java-based</li></ul>
Maestro (from Rice Univ)	<ul style="list-style-type: none"><li>•GPL</li><li>•Based on Java</li></ul>
NEC's Trema	<ul style="list-style-type: none"><li>•Open-source</li><li>•Written in C and Ruby</li><li>•Included test harness</li></ul>
Big Switch's Floodlight	<ul style="list-style-type: none"><li>•Apache license</li><li>•Java-based</li></ul>

# Sample Commercial Switches

Model	Virtualize	Notes	
HP Procurve 5400zl or 6600	1 OF instance per VLAN	<ul style="list-style-type: none"> <li>-LACP, VLAN and STP processing before OpenFlow</li> <li>-Wildcard rules or non-IP pkts processed in s/w</li> <li>-Header rewriting in s/w</li> <li>-CPU protects mgmt during loop</li> </ul>	
NEC IP8800	1 OF instance per VLAN	<ul style="list-style-type: none"> <li>-OpenFlow takes precedence</li> <li>-Most actions processed in hardware</li> <li>-MAC header rewriting in h/w</li> </ul>	
Brocade MLX routers	Multiple OF instance per switch	<ul style="list-style-type: none"> <li>-Hybrid OpenFlow switch with legacy protocols and OpenFlow coexisting</li> <li>-OpenFlow commands can override state created by legacy protocols</li> </ul>	
Pronto 3290 or 3780 with Pica8 or Indigo firmware	1 OF instance per switch	<ul style="list-style-type: none"> <li>-No legacy protocols (like VLAN, STP)</li> <li>-Most actions processed in hardware</li> <li>-MAC header rewriting in h/w</li> </ul>	

# Intro to RYU: OpenFlow Control



## Components:


- Provides interface for control and state and generates events
- Communicates using message passing

## Libraries:

- Functions called by components
- Ex: OF-Config, Netflow, sFlow, Netconf, OVSDB



# OpenFlow: base

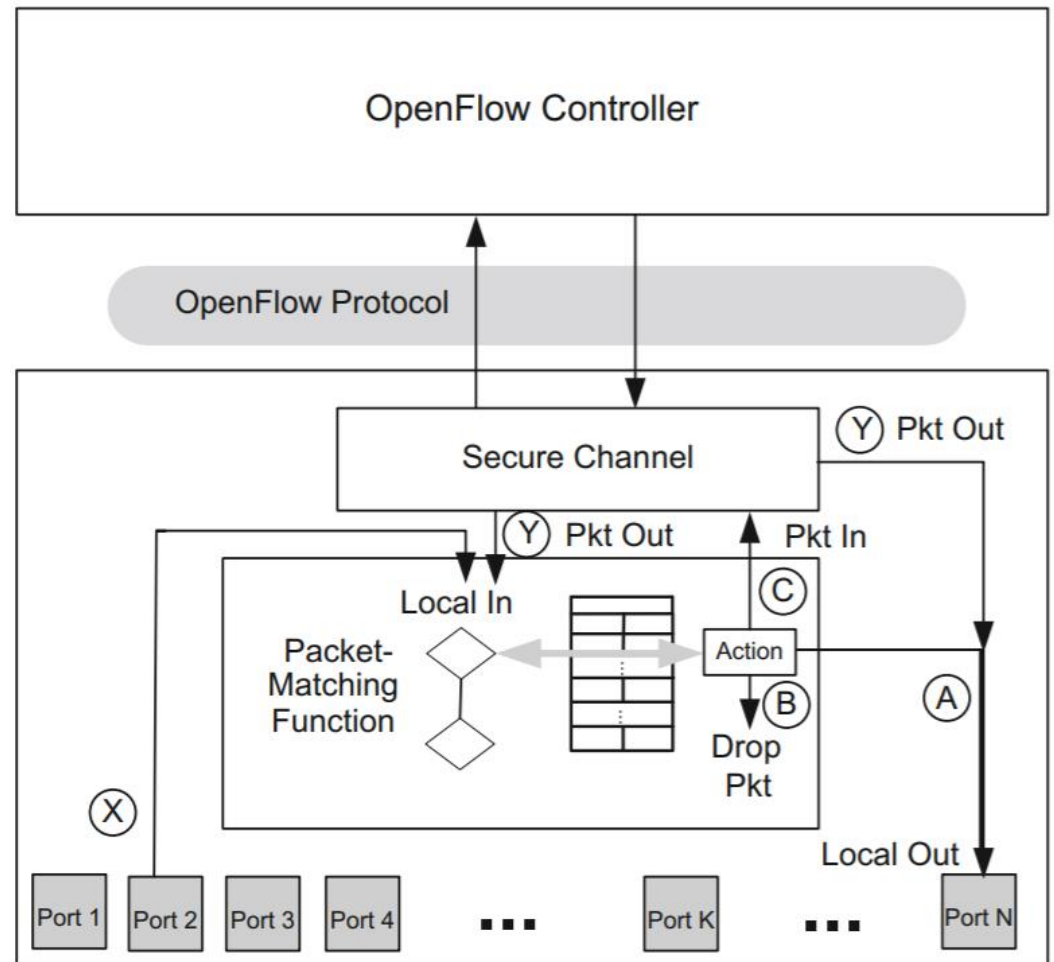


- Interface between switches and network controller

- A. Forward the packet out a local port, possibly modifying certain header fields first.

- B. Drop the packet

- C. Pass the packet to the controller



# OpenFlow: Flow Tables

- Header fields: match criteria to determine whether an incoming packet matches this entry
- Counters: track statistics relative to this flow, such as how many packets have been forwarded or dropped for this flow
- Actions fields prescribe what the switch should do with a packet matching this entry

Flow Entry 0		Flow Entry 1		...	Flow Entry F		...	Flow Entry M	
Header Fields	Inport 12 192.32.10.1, Port 1012	Header Fields	Inport * 209.*.*., Port *		Header Fields	Inport 2 192.32.20.1, Port 995		Header Fields	Inport 2 192.32.30.1, Port 995
Counters	val	Counters	val		Counters	val		Counters	val
Actions	val	Actions	val		Actions	val		Actions	val

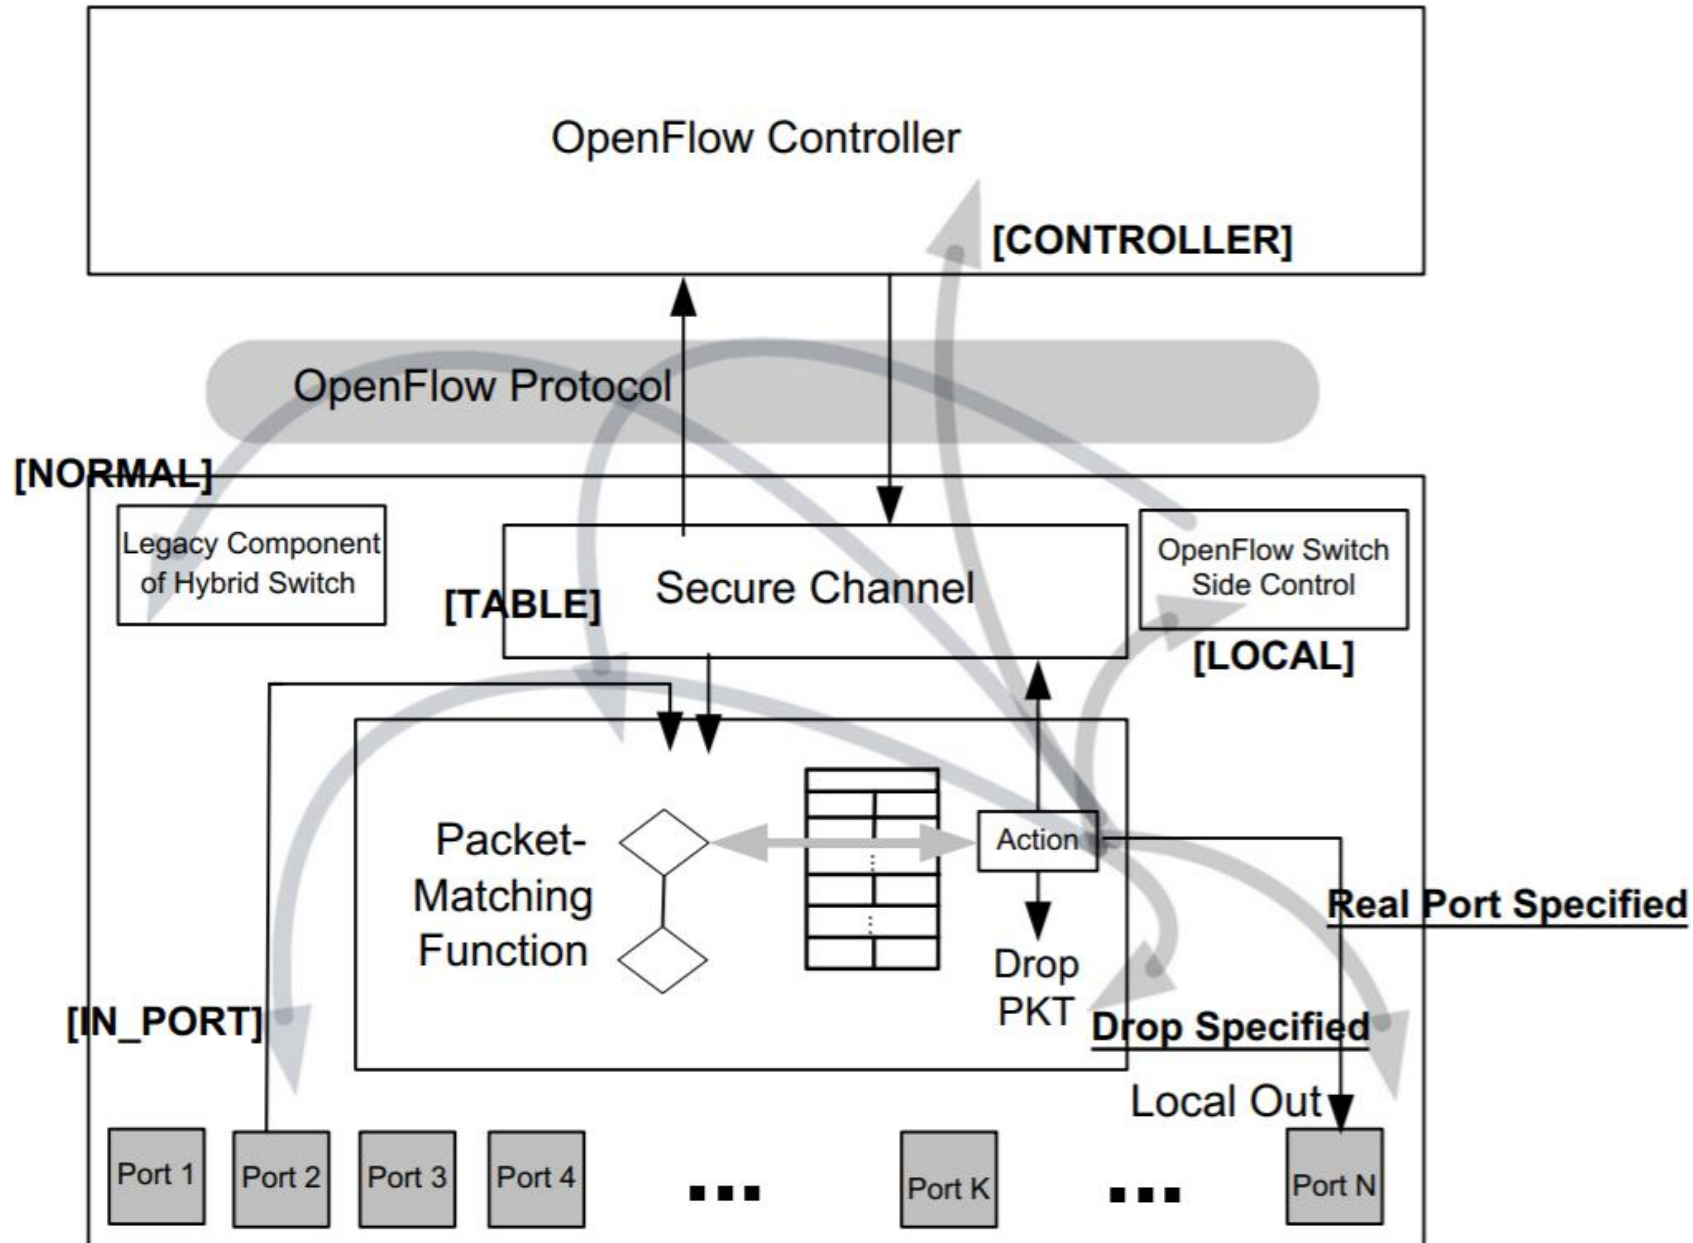
# OpenFlow: Packet Matching criteria

- Flow entries are processed in order: once a match is found, no further match attempts are made against that flow table
  - Switch input port
  - VLAN ID
  - VLAN priority
  - Ethernet source/destination address
  - Ethernet frame type
  - IP source/destination address
  - IP protocol
  - IP Type of Service (ToS) bits
  - TCP/UDP source/destination port
  - ...

# OpenFlow: Packet Forwarding Actions


- Local: packet needs to be processed by the local OpenFlow control software
- All/Flood: flood a packet out all ports on the switch except the input port
- Controller: forward to the OpenFlow controller
- In\_Port: forward the packet back out of the port on which it arrived (2 virtual machines on the same server to communicate)
- Table: packets that the controller sends to the switch, which arrive as part of the PACKET\_OUT message from the controller, which includes an action list
- Normal: forward to the legacy forwarding logic of the switch

# OpenFlow: Packet Forwarding



# OpenFlow Switching

Controller



Software  
Layer

OpenFlow Client

OpenFlow Table

MAC src	MAC dst	IP Src	IP Dst	TCP sport	TCP dport	Action
*	*	*	5.6.7.8	*	*	port 1

Hardware  
Layer

port 1

port 2

port 3

port 4

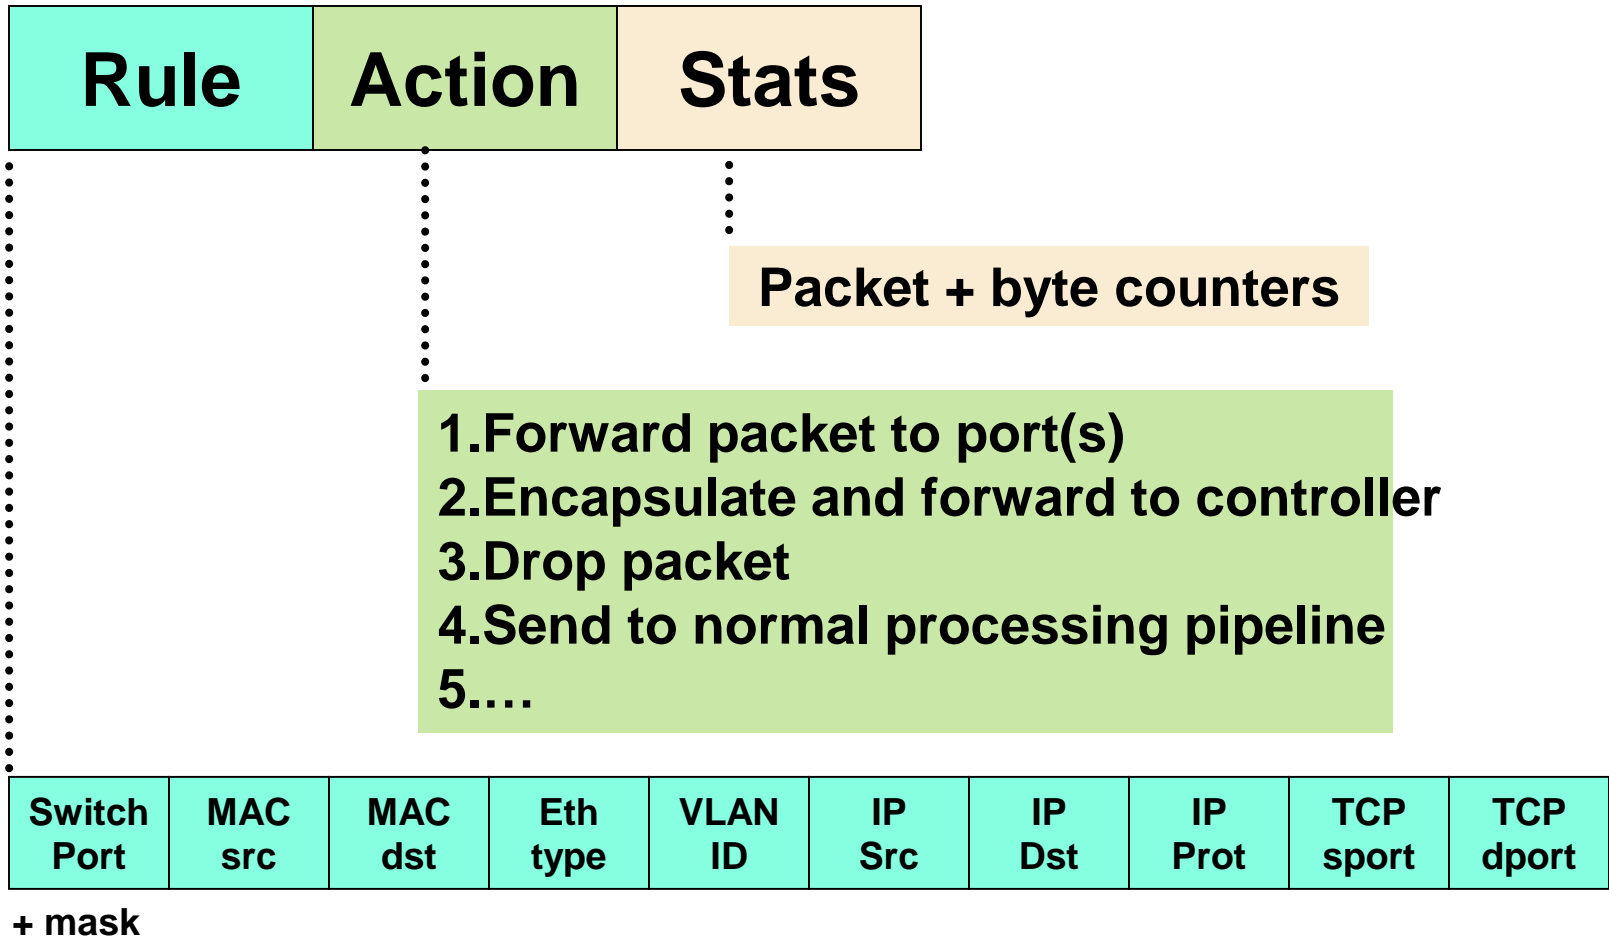


5.6.7.8



1.2.3.4

# OpenFlow Table Entry



# OpenFlow Examples

## Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f:...	*	*	*	*	*	*	*	port6

## Routing

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	5.6.7.8	*	*	*	port6

## Firewall

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	*	*	*	22	drop



# OpenFlow Examples

## Flow Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
port3	00:20..	00:1f..	0800	vlan1	1.2.3.4	5.6.7.8	4	17264	80	port6

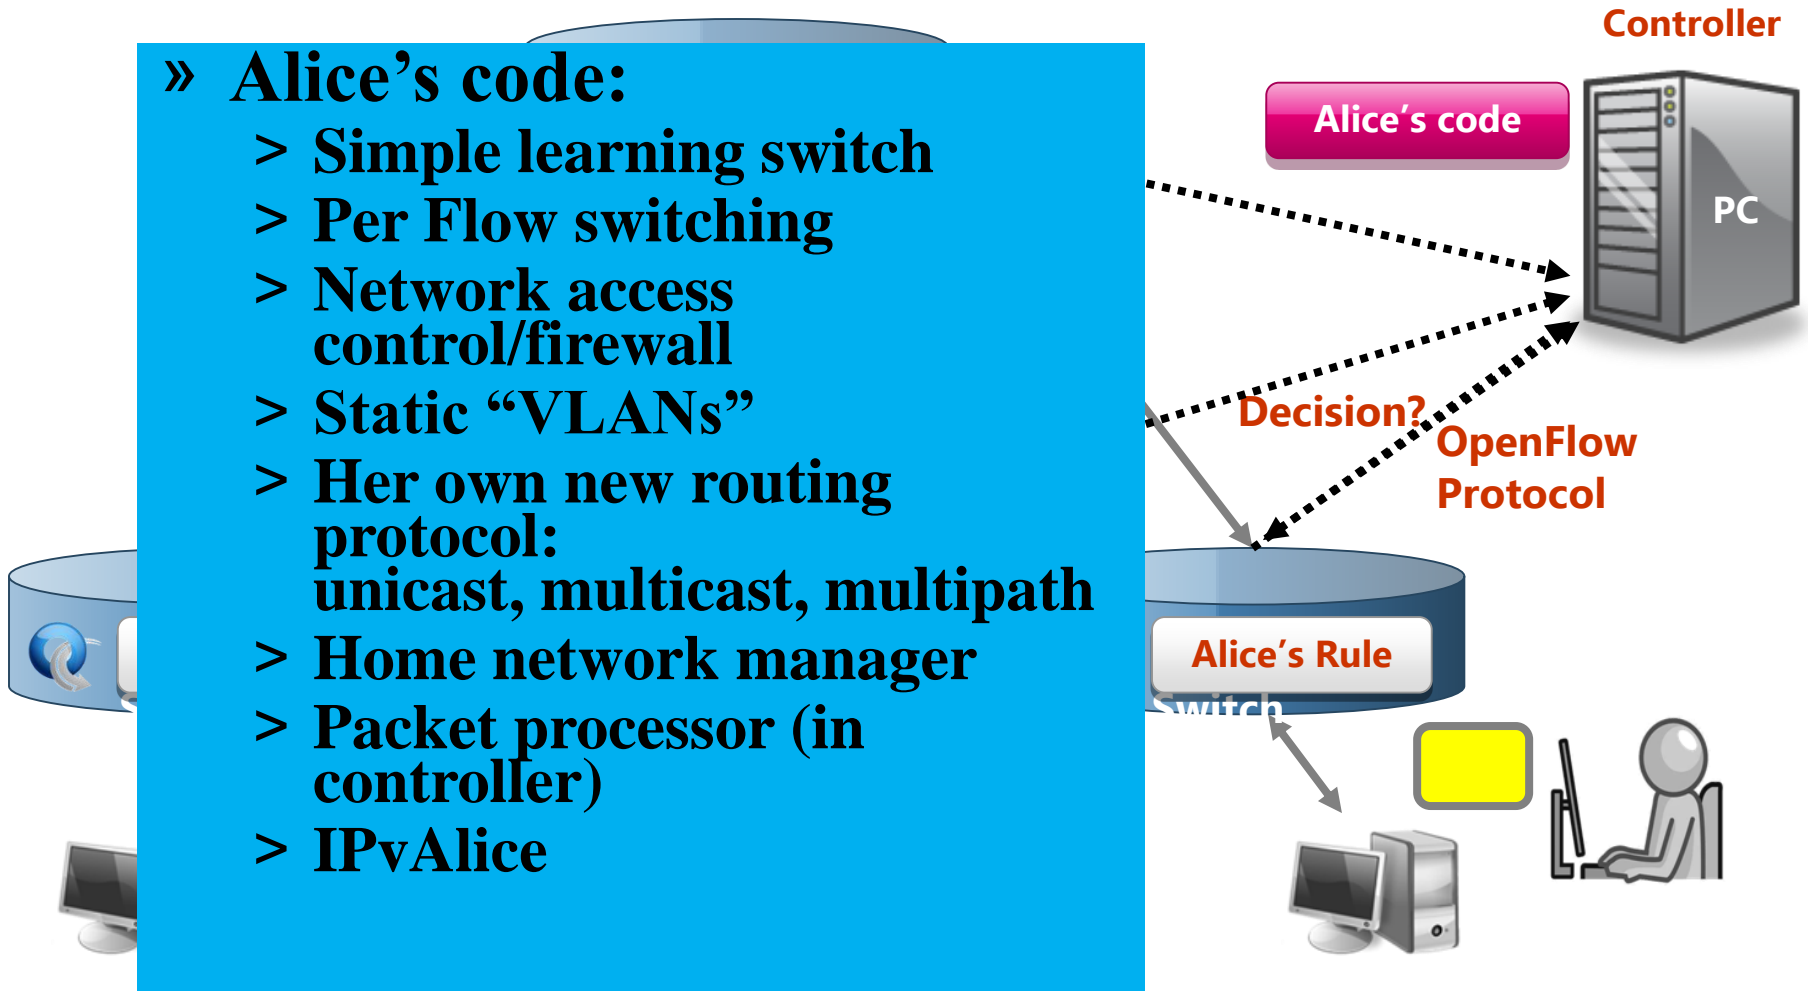
## VLAN Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f..	*	vlan1	*	*	*	*	*	port6, port7, port9

# OpenFlow Usage

## » Alice's code:

- > Simple learning switch
- > Per Flow switching
- > Network access control/firewall
- > Static "VLANs"
- > Her own new routing protocol:  
unicast, multicast, multipath
- > Home network manager
- > Packet processor (in controller)
- > IPvAlice



# OpenFlow Rules in Different Switches


- Access control rules placement
- Firewall policy is transformed into:
  - list of rules R1, R2, R3 used to block matching packets
  - two default rules RD1, RD2 for forwarding non-matching packets towards endpoint E
  - Rules distributed on several switches to ensure that flows F1 and F2 pass through all rules R1, R2, R3 to satisfy the policy.



Block traffic to **10.0.0.1** or to port **22** or from **10.0.0.2**

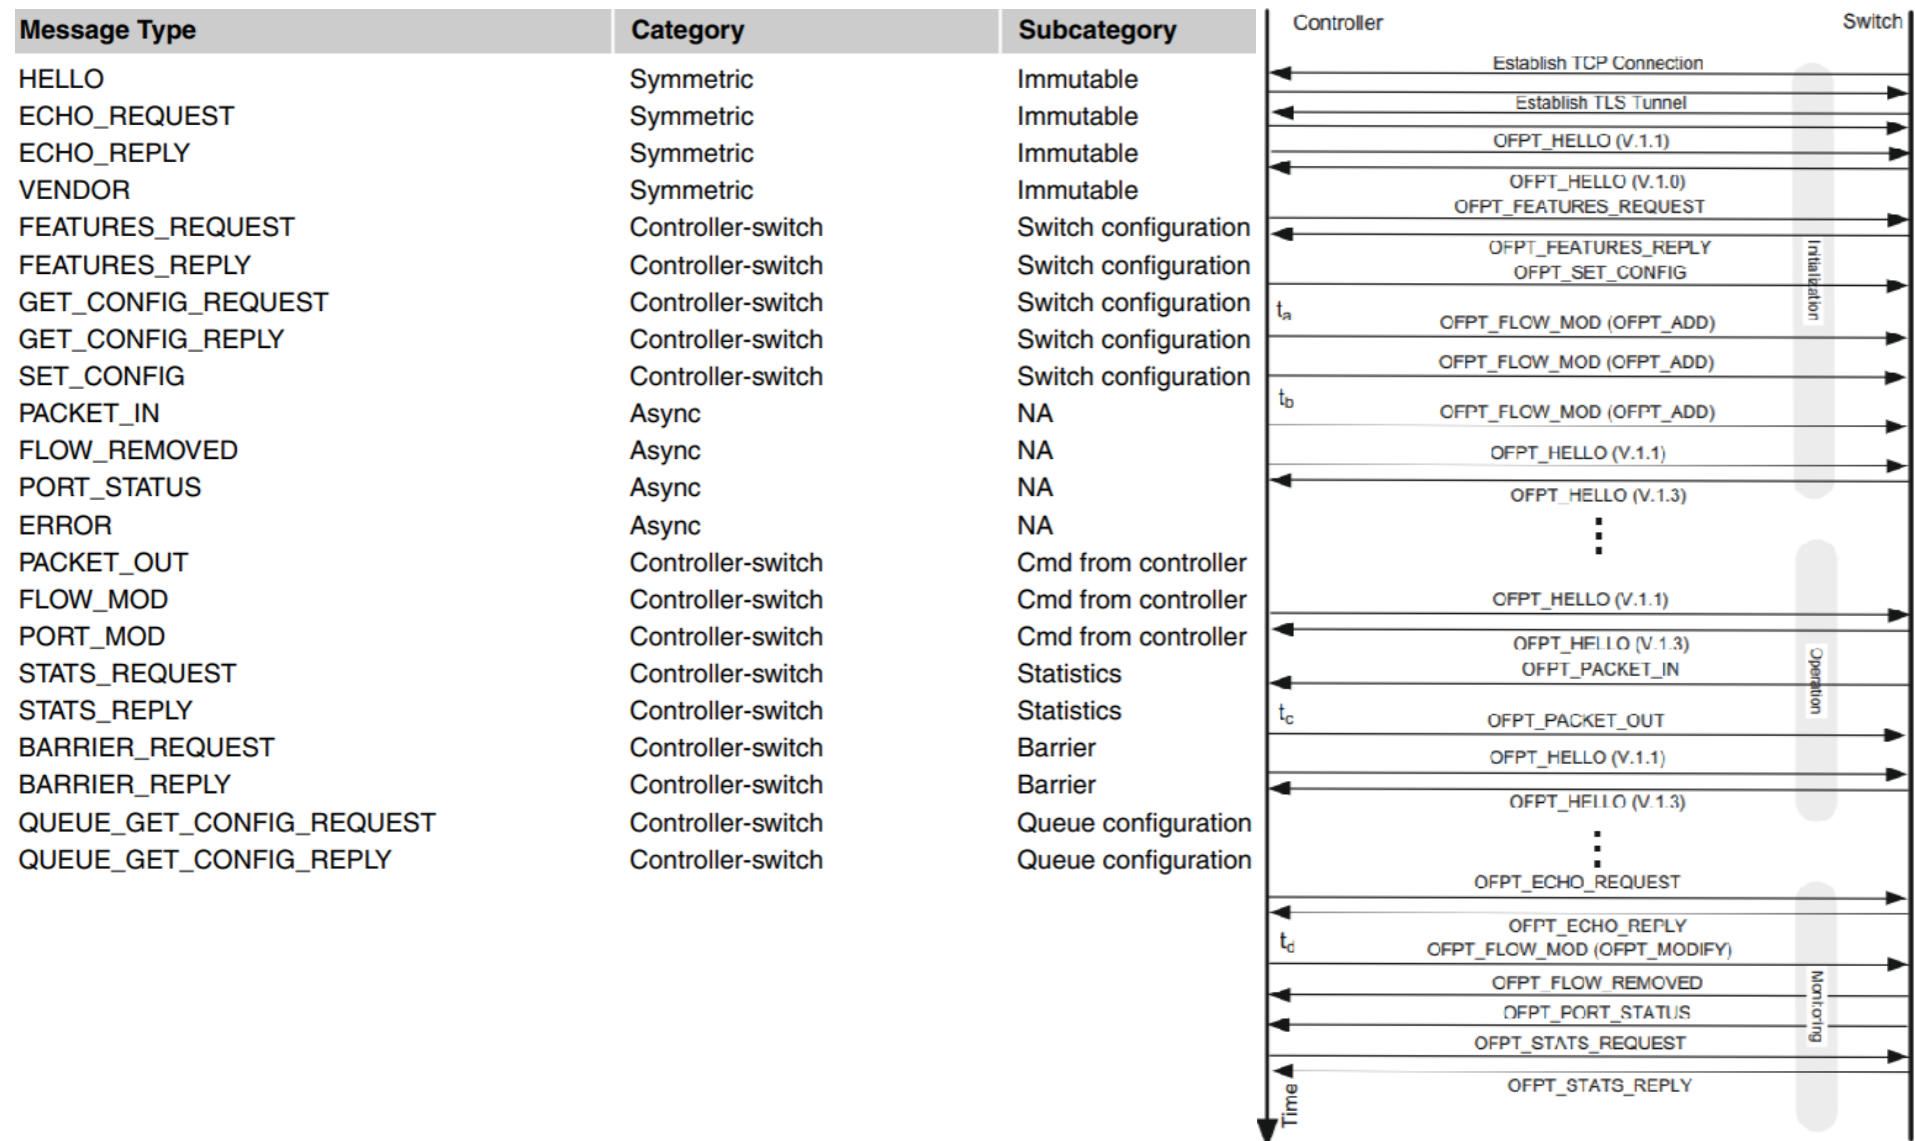


Rule	Match	Action
R1	dst_ip=10.0.0.1	Drop
R2	dst_port=22	Drop
R3	src_ip=10.0.0.2	Drop
RD1	*	To Switch 2
RD2	*	E1



Switch Memory Capacity = 3

# OpenFlow: Messages (examples)

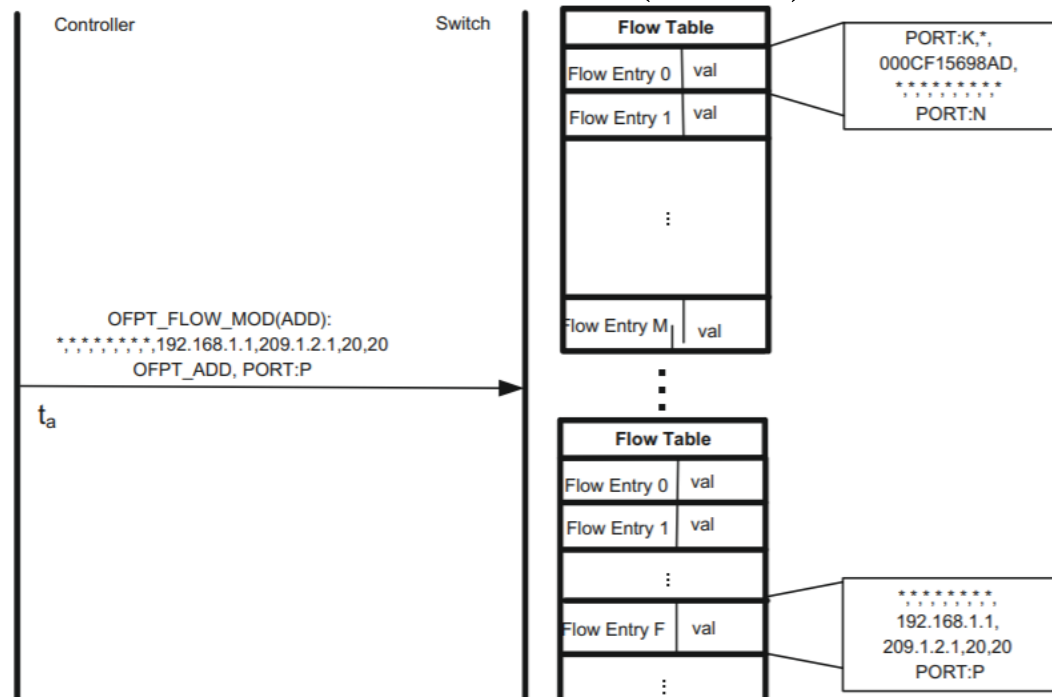


# OpenFlow: Controller

## Programming Flow Table

- Entry 0: all Ethernet frames entering the switch on input port K with a destination Ethernet address of 0x000CF15698AD should be output on output port N
- At time  $t_a$ , the controller sends a FLOW\_MOD (ADD)

command to the switch, adding a flow for packets with source IP addresses 192.168.1.1 and destination 209.1.2.1, source TCP port 20, and destination port 20. All other match fields have been wildcarded

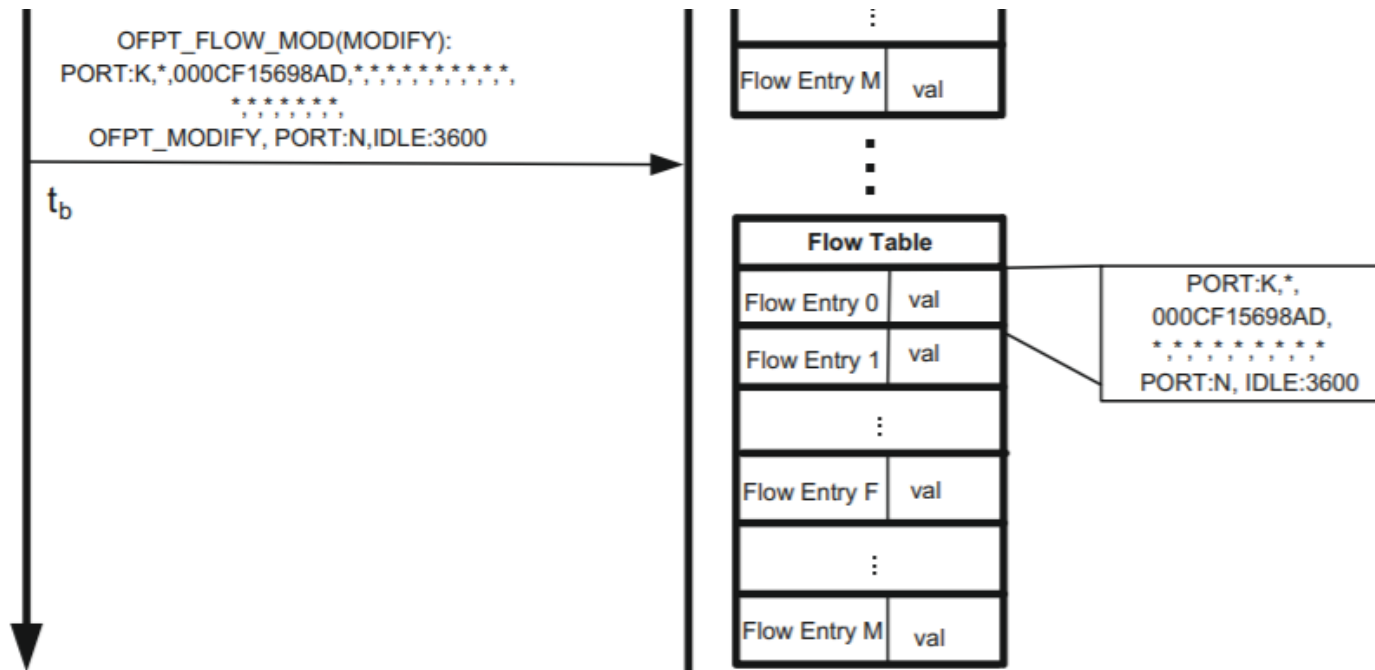


- The output port is specified as P

# OpenFlow: Controller


## Programming Flow Table

- FLOW\_MOD (MODIFY) command for flow entry zero
- One-hour (3600-second) idle time on that flow
  - After that number of seconds of inactivity on that flow, the flow should be deleted

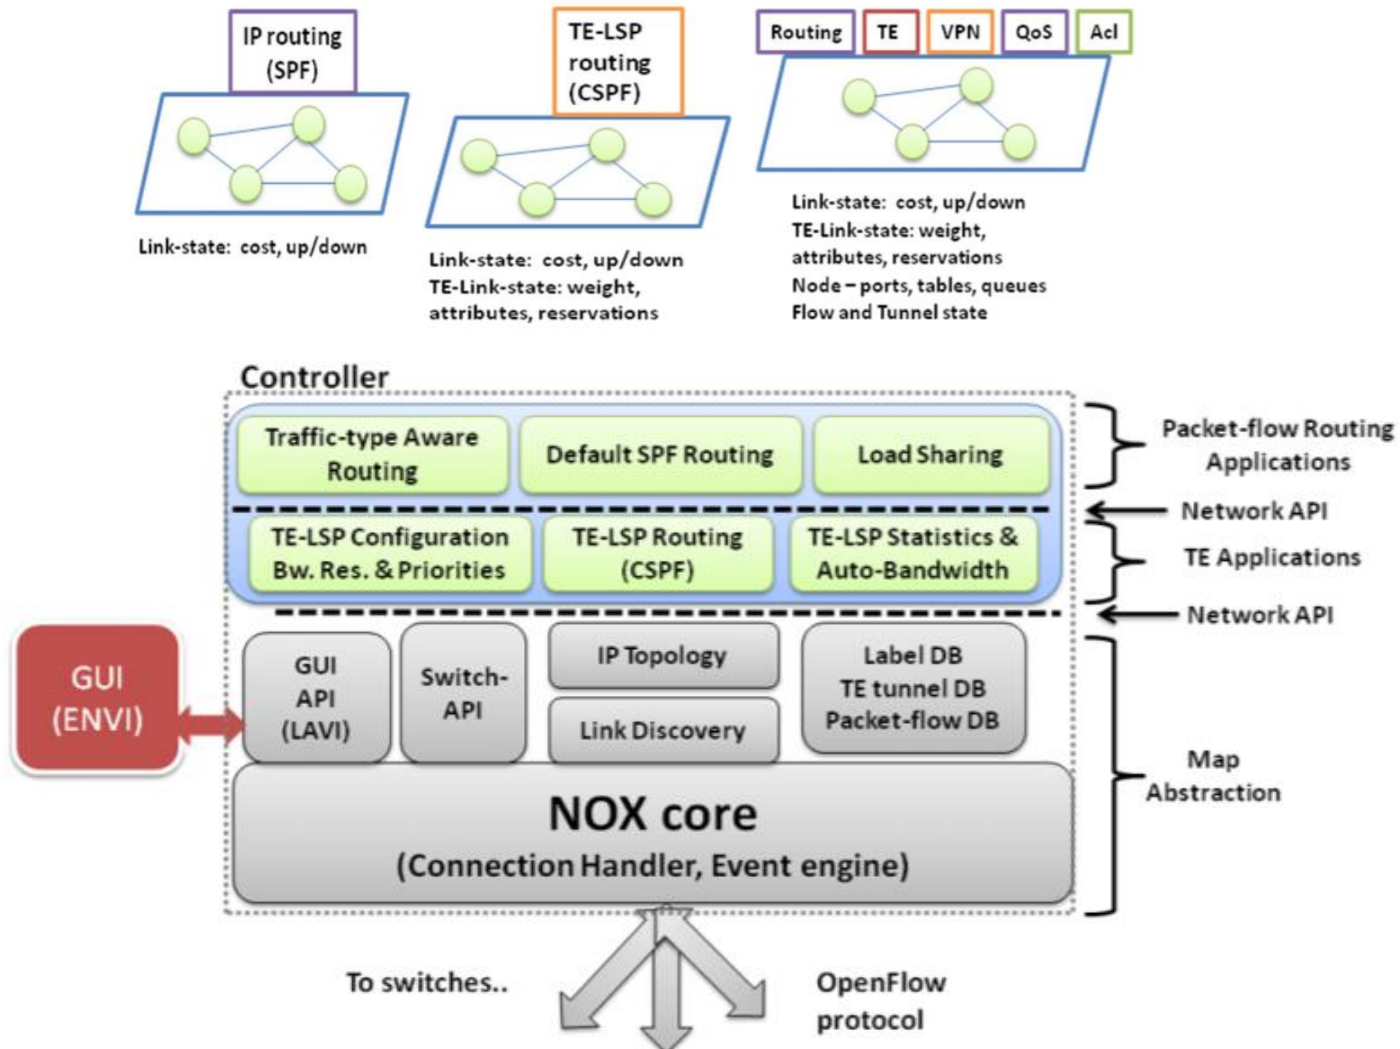


# OpenFlow: Basic Packet Forwarding

- Packet arriving through port 2 with source IPv4 address of 192.168.1.1 and destination IPv4 address of 209.1.2.1
- Packet-matching function scans the flow table
- Finds a match in flow entry F → matching packet should be forwarded out port P



# OpenFlow: MPLS



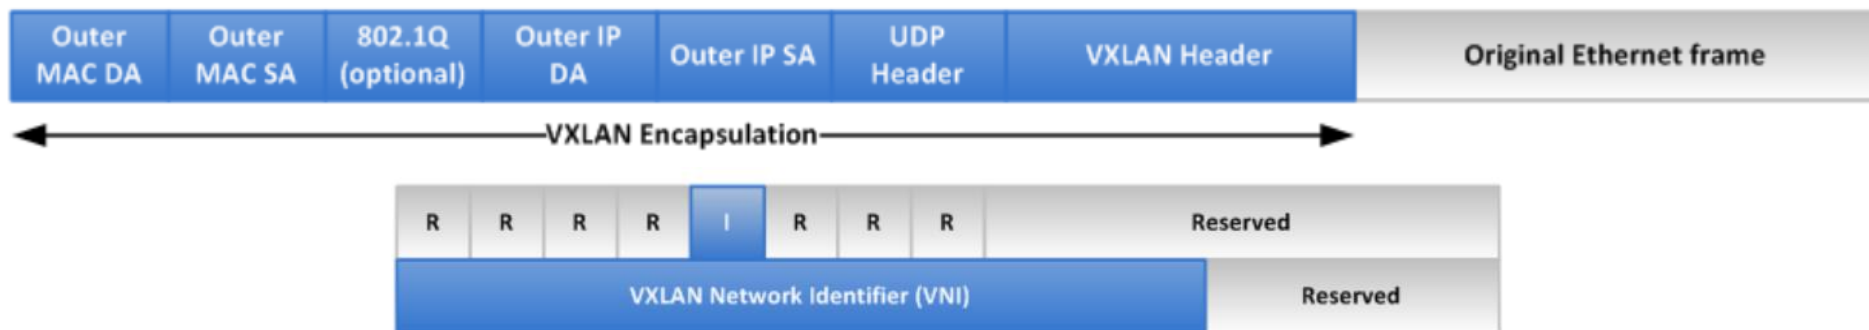


# Virtual Traffic and Tunnels

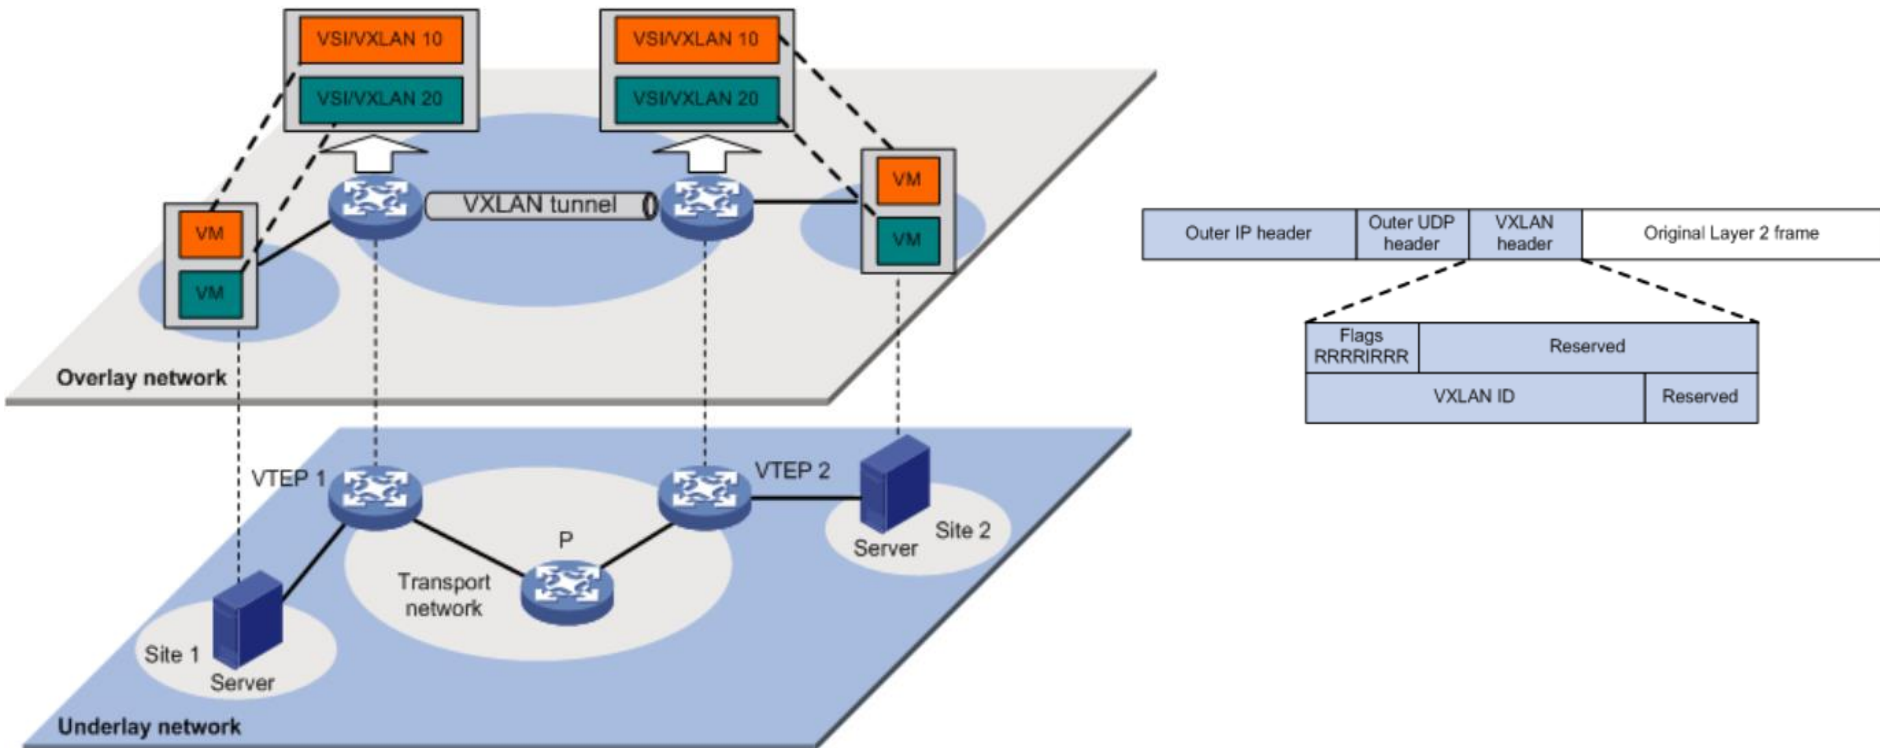
- Virtual network traffic runs above the physical network infrastructure
  - Endpoints are unaware of the details of the physical topology, the way routing occurs, or other basic network functions
  - Virtual networks can be controlled entirely by the devices at the edge of the network → Tunneling
  - When a packet enters the edge of the virtual network at the source, the networking device will encapsulate it within another frame
    - Edge of the virtual network: virtual tunnel endpoint (VTEP)
  - It is then sent, through information programmed by the controller, to the destination's VTEP
  - This VTEP decapsulates the packet and forwards it to the destination host
  - MAC-in-IP tunneling: VXLAN (Virtual eXtensible LAN), NVGRE, STT
  - Multiple overlay networks can exist independently and simultaneously over the same physical network

# VXLAN: Virtual eXtensible LAN

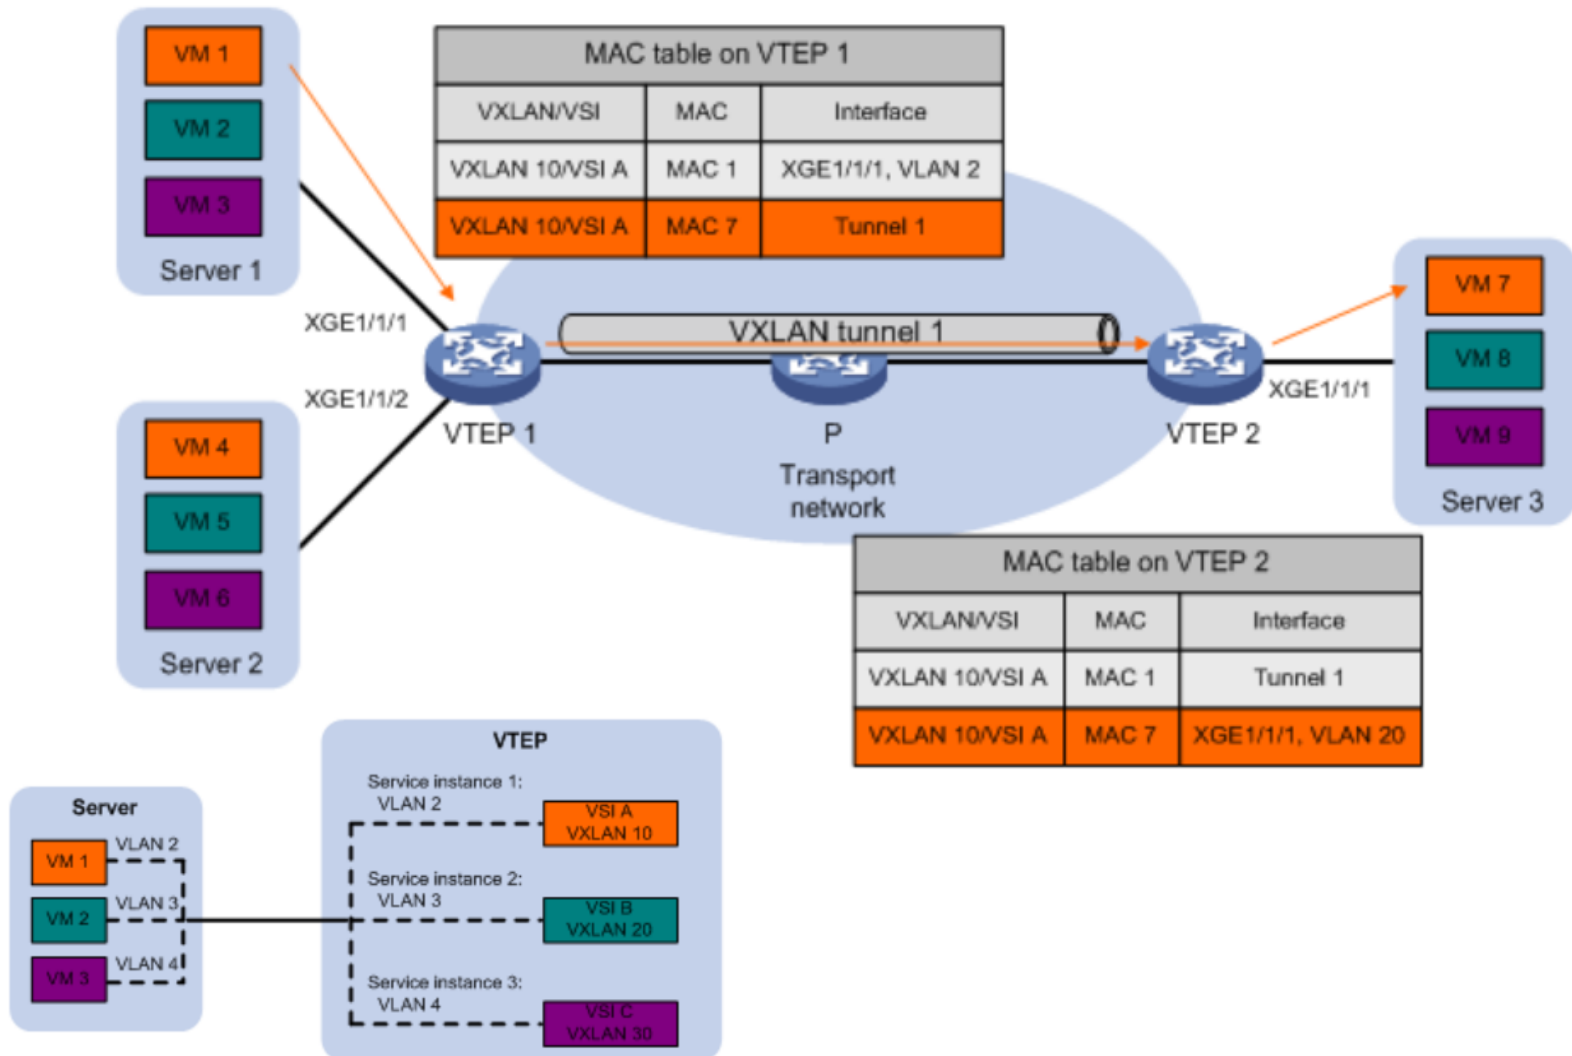
- Same services as VLANs, in a virtualized environment
- Also provide a means to stretch L2 network over a L3 network
- VXLAN ID (called VXLAN Network Identifier or VNI) is 24-bits long compared to 12-bits of VLAN ID (over 16 million unique IDs)
- VTEPs connect an access switch (currently virtual switch) to the IP network
- The function of VTEP is to encapsulate the virtual traffic within an IP header to send across an IP network



# VXLAN Example




# VXLAN Example




# Centralized vs Distributed Control

## Centralized Control




## Distributed Control

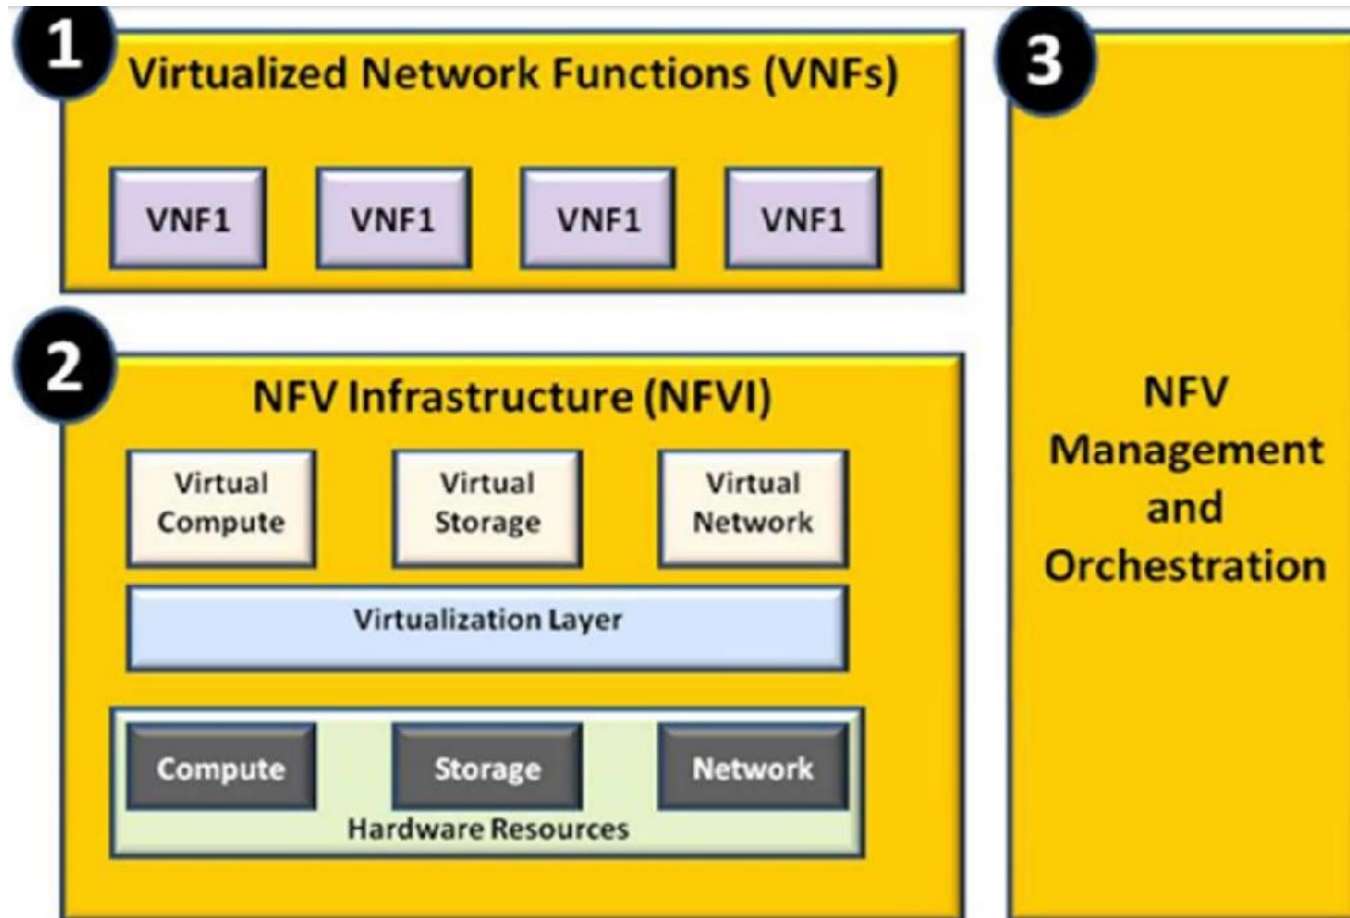


# Network Functions Virtualization

- **Management & orchestration**
  - **infrastructure management standards**
  - **multi-level identity standard**
  - **resource description language**



# NFV Architecture



# NFV Examples

